

---

# **Data Science Project**

## **Scala Report**

**Prepared by Group 6**

### **Team members**

Bhavana Nare

Vishnu Priya Buggineni

Yaswanth Kumar Palli

Vamshi Krishna B

We have worked on all 6 datasets using both **python** and **scala**.

All the libraries we used in python are mentioned below:

Numpy, Pandas, Matplotlib, Seaborn, Mlxtend, Gplearn, Math, and sklearn

#### **Air Quality Dataset:**

We have taken this dataset from UCI repository(We used AirQualityUCI.csv for scala). There are 9358 occurrences of hourly averaged responses from an array of 5 metal oxide chemical sensors integrated in an Air Quality Chemical Multisensor Device in this dataset. This dataset contains 12 attributes, the information about each attribute is mentioned below:

Date- format (DD/MM/YYYY)

Time- format (HH.MM.SS)

CO(GT)- True hourly averaged concentration CO

PT08.S1(CO)- (tin oxide) hourly averaged sensor response

NMHC(GT)- True hourly averaged overall Non Metanic HydroCarbons concentration

C6H6(GT)- True hourly averaged Benzene concentration

PT08.S2(NMHC)- (titania) hourly averaged sensor response

NOx(GT)- True hourly averaged NOx concentration

PT08.S3(NOx)- (tungsten oxide) hourly averaged sensor response

NO2(GT)- True hourly averaged NO2 concentration—**Response Variable**

PT08.S4(NO2)- (tungsten oxide) hourly averaged sensor response

PT08.S5(O3) -(indium oxide) hourly averaged sensor response

T- Temperature in °C

RH- Relative Humidity (%)

AH- Absolute Humidity

#### **Scala report:**

First we imported all the necessary modules.We used matrixD.load() to load the dataset.

The date and time are excluded from the airquality dataset. The airquality dataset contains negative values, So we used imputation to replace the -200 values with mean. The code we used for imputation is mentioned below:

ession1.scala •


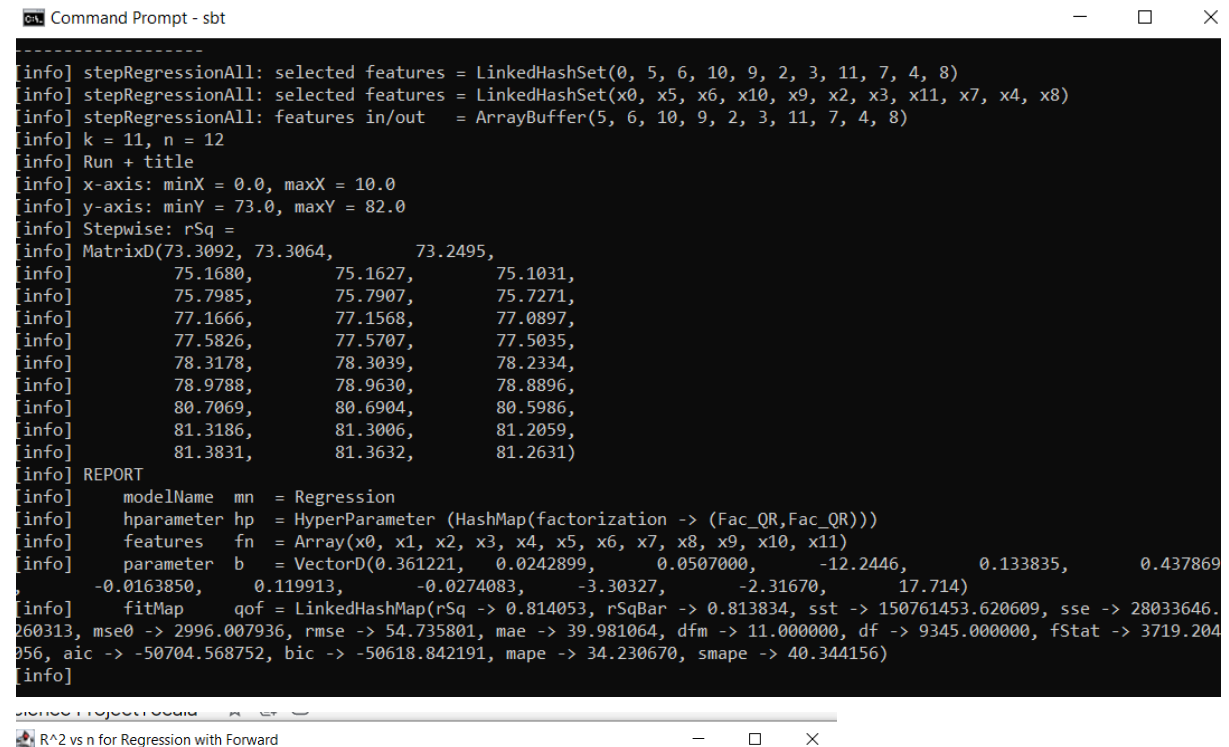
main > scala > scalation > modeling > regression1.scala

```
@main def RegressionTest12 (): Unit =
{
  val airqual = MatrixD.load("AirQualityUCI.csv",1)
  println(airqual.dim)
  println(airqual.dim2)
  val n = airqual.dim2 - 1
  val (x, y) = (airqual.not(?, n), airqual(?, n))
  for (i <- (0 until airqual.dim2) )
  {
    val z = airqual(?, i)
    val x2= x.copy
    val x3 = x.copy
    //var iv = ( -200.000,-200.000)
    val iv = ImputeMean.impute (z)
    println(iv._2)
    println(z.length)

    for (j <- (0 until z.length) ){
      if (airqual(j, i) == -200.000){
        airqual(j, i) = iv._2
      }
    }
  }

  airqual.write("test.csv")
}
```

## Linear Regression:




**R2: 0.814053**

**Adj R2: 0.813834**

## Ridge Regression:

```
Command Prompt - sbt
93.000, 193.000, 185.000, 147.000, 117.000, 112.000, 101.000, 93.0000, 6
1.0000, 61.0000, 54.0000, 64.0000, 80.0000, 133.000, 170.000, 139.000, 1
24.000, 118.000, 108.000, 100.000, 91.0000, 89.0000, 102.000, 114.000, 1
24.000, 136.000, 118.000, 93.0000, 72.0000, 69.0000, 61.0000, 50.0000, 3
5.0000, 29.0000, 41.0000, 66.0000, 89.0000, 150.000, 150.000, 127.000, 1
20.000, 123.000, 115.000, 122.000, 111.000, 119.000, 137.000, 147.000, 1
64.000, 160.000, 134.000, 110.000, 99.0000, 127.000, 174.000, 127.000, 8
1.0000, 70.0000, 49.0000, 44.0000, 107.000, 141.000, 121.000, 103.000, 1
92.000, 95.0000, 93.0000, 95.0000, 80.0000, 81.0000, 106.000, 113.000, 1
7.000, 124.000, 122.000, 87.0000, 79.0000, 93.0000, 102.000, 86.0000, 7
7.0000, 88.0000, 64.0000, 46.0000, 55.0000, 84.0000, 89.0000, 118.000, 1
3.000, 115.000, 124.000, 86.0000, 87.0000, 100.000, 132.000, 156.000, 1
8.000, 181.000, 187.000, 158.000, 153.000, 128.000, 93.0000, 58.0000, 5
5.0000, 51.0000, 43.0000, 53.0000, 93.0000, 155.000, 174.000, 187.000, 1
0.000, 179.000, 175.000, 156.000, 168.000)
[info] -----
[info] | Optimize lambda |
[info] -----
[info] findLambda2 = -0.0
[info] REPORT
[info] modelName mn = RidgeRegression
[info] hparameter hp = HyperParameter (HashMap(lambda -> (1.0,0.01), factorization -> (Fac_Cholesky,Fac_Cholesky))
[info] features fn = Array(x0, x1, x2, x3, x4, x5, x6, x7, x8, x9, x10, x11)
[info] parameter b = VectorD(0.361227, 0.0242961, 0.0507017, -12.2433, 0.133810, 0.437866
-0.0163796, 0.119900, -0.0274074, -3.30269, -2.31651, 17.712)
[info] fitMap qof = LinkedHashMap(rSq -> 0.814053, rSqBar -> 0.813814, sst -> 150761453.620609, sse -> 28033646.
15129, mse0 -> 2996.007942, rmse -> 54.735801, mae -> 39.981120, dfm -> 12.000000, df -> 9344.000000, fStat -> 3408.905
654, aic -> -50702.568761, bic -> -50609.698321, mape -> 34.230603, smape -> 40.343996)
[info]
```

Reset Plot




R2: 0.814053

Adj R2: 0.813814

## Lasso Regression:

```
Command Prompt - sul
149.99, 64.4204, 53.6079, 1.00000, 9355.00, 27048.0, -10444.4, -10433.3, 4
9.6001, 59.490)
[info] DEBUG @ Predictor.crossValidate: fold 3: test set size = 1871
[info] DEBUG @ Predictor.validate: test set size = 1871
[info] DEBUG @ LassoRegression.train: LassoAdmm estimates parameter b = VectorD(0.413841, 0.36228)
[info] DEBUG @ Predictor.validate: parameters b = VectorD(0.413841, 0.36228)
[info] DEBUG @ Predictor.crossValidate: fold 3: qof = VectorD(0.711937, 0.711906, 2.94996e+07, 8.49774e+06, 4
541.82, 67.3930, 55.5291, 1.00000, 9355.00, 23120.5, -10531.8, -10520.7, 5
0.6385, 62.118)
[info] DEBUG @ Predictor.crossValidate: fold 4: test set size = 1871
[info] DEBUG @ Predictor.validate: test set size = 1871
[info] DEBUG @ LassoRegression.train: LassoAdmm estimates parameter b = VectorD(0.399677, 0.36656)
[info] DEBUG @ Predictor.validate: parameters b = VectorD(0.399677, 0.36656)
[info] DEBUG @ Predictor.crossValidate: fold 4: qof = VectorD(0.732676, 0.732647, 3.02237e+07, 8.07953e+06, 4
318.29, 65.7137, 53.9803, 1.00000, 9355.00, 25640.0, -10481.9, -10470.9, 4
8.8367, 58.883)
[info] -----
[info] | backwardElimAll: (1 = 10) REMOVE variable (7, x7) => cols = LinkedHashSet(0, 5) @ 0.733063524968985 |
[info] -----
[info] REPORT
[info] modelName mn = LassoRegression
[info] hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01)))
[info] features fn = Array(x0, x1, x2, x3, x4, x5, x6, x7, x8, x9, x10, x11)
[info] parameter b = VectorD(0.361221, 0.0242899, 0.0507000, -12.2446, 0.133835, 0.437869
, -0.0163850, 0.119913, -0.0274083, -3.30327, -2.31670, 17.714)
[info] fitMap qof = LinkedHashMap(rSq -> 0.814053, rSqBar -> 0.813834, sst -> 150761453.620609, sse -> 28033646.
260312, mse0 -> 2996.007936, rmse -> 54.735801, mae -> 39.981064, dfm -> 11.000000, df -> 9345.000000, fStat -> 3719.204
056, aic -> -50704.568752, bic -> -50618.842191, mape -> 34.230670, smape -> 40.344156)
[info]
```




R2: 0.814053

Adj R2: 0.813834

## Quadratic Regression:

```
Command Prompt - sbt

[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | 0.576 | 0.638 | 0.608 | 0.022 | 0.028 |
[info] | rSqBar | 5 | 0.576 | 0.638 | 0.608 | 0.022 | 0.028 |
[info] | sst | 5 | 82832113.983 | 87920724.853 | 84690782.380 | 2002528.300 | 2486954.372 |
[info] | sse | 5 | 30805228.308 | 35319171.311 | 33205180.373 | 1753807.782 | 2178066.563 |
[info] | mse0 | 5 | 16464.580 | 18877.163 | 17747.290 | 937.364 | 1164.119 |
[info] | rmse | 5 | 128.314 | 137.394 | 133.182 | 3.530 | 4.384 |
[info] | mae | 5 | 94.416 | 100.455 | 97.315 | 2.297 | 2.853 |
[info] | dfm | 5 | 2.000 | 2.000 | 2.000 | 0.000 | 0.000 |
[info] | df | 5 | 9354.000 | 9354.000 | 9354.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 6359.958 | 8242.418 | 7278.074 | 679.657 | 844.071 |
[info] | aic | 5 | -11861.350 | -11734.099 | -11801.755 | 49.441 | 61.401 |
[info] | bic | 5 | -11844.747 | -11717.496 | -11785.152 | 49.441 | 61.401 |
[info] | mape | 5 | 8.750 | 9.294 | 9.033 | 0.258 | 0.320 |
[info] | smape | 5 | 8.554 | 9.100 | 8.820 | 0.228 | 0.283 |
[info] -----
[info] REPORT
[info] modelName mn = PolyRegression
[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac_Cholesky,Fac_Cholesky)))
[info] features fn = Array(x0, x1, x2)
[info] parameter b = VectorD(784.874, 164.252, -6.1798)
[info] fitMap qof = LinkedHashMap(rSq -> 0.608218, rSqBar -> 0.608134, sst -> 423642445.749750, sse -> 165975457.961019, mse0 -> 17738.106013, rmse -> 133.184481, mae -> 97.271774, dfm -> 2.000000, df -> 9354.000000, fStat -> 7260.763228, aic -> -59042.974549, bic -> -59021.542909, mape -> 9.024178, smape -> 8.817180)
[info]
```



R2: 0.608218

Adj R2: 0.608134


## Symbolic Regression:

Command Prompt - sbt

```
[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac_QR, Fac_QR)))
[info] features fn = Array(CO(GT), PT08.S1(CO), NMHC(GT), C6H6(GT), PT08.S2(NMHC), NOx(GT), PT08.S3(NOx), PT08.S4(NO2), PT08.S5(O3), T, RH, AH, PT08.S1(CO)_CO(GT), NMHC(GT)_CO(GT), NMHC(GT)_PT08.S1(CO), C6H6(GT)_CO(GT), C6H6(GT)_PT08.S1(CO), C6H6(GT)_NMHC(GT), PT08.S2(NMHC)_CO(GT), PT08.S2(NMHC)_PT08.S1(CO), PT08.S2(NMHC)_NMHC(GT), PT08.S2(NMHC)_C6H6(GT), NOx(GT)_CO(GT), NOx(GT)_PT08.S1(CO), NOx(GT)_NMHC(GT), NOx(GT)_C6H6(GT), NOx(GT)_PT08.S2(NMHC), PT08.S3(NOx)_CO(GT), PT08.S3(NOx)_PT08.S1(CO), PT08.S3(NOx)_NMHC(GT), PT08.S3(NOx)_C6H6(GT), PT08.S3(NOx)_PT08.S2(NMHC), PT08.S3(NOx)_NOx(GT), PT08.S4(NO2)_CO(GT), PT08.S4(NO2)_PT08.S1(CO), PT08.S4(NO2)_NMHC(GT), PT08.S4(NO2)_C6H6(GT), PT08.S4(NO2)_PT08.S2(NMHC), PT08.S4(NO2)_NOx(GT), PT08.S4(NO2)_PT08.S3(NOx), PT08.S5(O3)_CO(GT), PT08.S5(O3)_PT08.S1(CO), PT08.S5(O3)_NMHC(GT), PT08.S5(O3)_C6H6(GT), PT08.S5(O3)_PT08.S2(NMHC), PT08.S5(O3)_NOx(GT), PT08.S5(O3)_PT08.S3(NOx), PT08.S5(O3)_PT08.S4(NO2), T_CO(GT), T_PT08.S1(CO), T_NMHC(GT), T_C6H6(GT), T_PT08.S2(NMHC), T_NOx(GT), T_PT08.S3(NOx), T_PT08.S4(NO2), T_PT08.S5(O3), RH_CO(GT), RH_PT08.S1(CO), RH_NMHC(GT), RH_C6H6(GT), RH_PT08.S2(NMHC), RH_NOx(GT), RH_PT08.S3(NOx), RH_PT08.S4(NO2), RH_PT08.S5(O3), RH_T, AH_CO(GT), AH_PT08.S1(CO), AH_NMHC(GT), AH_C6H6(GT), AH_PT08.S2(NMHC), AH_NOx(GT), AH_PT08.S3(NOx), AH_PT08.S4(NO2), AH_PT08.S5(O3), AH_T, AH_RH)
[info] parameter b = VectorD(-0.976034, -0.135022, 0.462958, -28.6695, 0.102849, 1.96163, -0.0412362, -0.105958, 0.0463079, 1.02112, 2.96906, -180.036, 0.000858809, -0.000277666, -0.000101399, -0.0143979, -0.000115742, 0.0117626, 0.000463475, 0.000450867, -0.000200400, 0.00754315, -0.000986516, -0.000361759, 0.000254492, 0.0280100, -0.00152005, 0.000147738, 0.000157151, -0.000162943, 0.00692277, -0.000342181, -5.68531e-05, 0.000307364, -0.000300042, -0.000222607, -0.0194473, 0.000616826, 0.000244363, -4.10545e-06, -0.000178948, 0.000131338, -4.27895e-05, 0.0110176, -0.000412535, 9.07251e-06, 3.73794e-05, 1.87043e-05, -0.0170942, -0.00165674, 0.00775525, 0.379995, -0.00991272, 0.00261205, 0.00422417, 0.00297687, -5.06375e-05, -0.00565141, 0.00162702, 0.00307199, 0.173256, -0.00440764, -0.00201795, 0.000933011, -0.000643902, -0.000815811, 0.0683510, 0.0273173, -0.0535103, -0.0195080, -3.89541, -0.0526628, -0.0184314, -0.0601777, 0.100579, 0.0312307, 2.67829, -0.39155)
[info] fitMap qof = LinkedHashMap(rSq -> 0.952488, rSqBar -> 0.952093, sst -> 150761453.620609, sse -> 7163045.514921, mse0 -> 765.528002, rmse -> 27.668177, mae -> 19.072004, dfm -> 77.000000, df -> 9279.000000, fStat -> 2415.807717, aic -> -44188.894925, bic -> -43631.672284, mape -> 17.038205, smape -> 17.521613)
[info]
[info] Run + title
```

SymbolicRegressionX: y actual, predicted

Reset Plot



R2: 0.952488

Adj R2: 0.952093



## Symbolic Ridge Regression:

```
-0.07, -8.84040e-07, -4.47320e-05, 1.05508e-06, 3.31231e-07, -3.24656e-07, -2.23063e-07, -5.16524e-06, 6.60529e-07, 2.18323e-08, -6.23063e-07, 5.80493e-07, 4.88642e-07, 6.45662e-05, -2.38438e-06, -4.18862e-06, -9.92365e-07, 5.04494e-07, -2.14042e-07, -9.32604e-06, 3.66260e-08, 5.82433e-07, 6.78548e-07, -7.12718e-06, 5.53967e-08, 4.55389e-07, 2.27909e-07, 1.64509e-07, -8.30708e-06, 6.10063e-08, 1.67320e-07, 3.52666e-07, -6.32627e-07, -2.21246e-07, 1.83245e-05, 2.80966e-08, -5.14840e-07, -1.52283e-08, -7.21392e-05, -0.000139480, 1.70203e-05, 0.000214924, -0.000545448, 0.00218834, -6.84406e-05, 1.94100e-05, 4.67947e-05, 0.00139561, 5.99093e-05, 2.31865e-05, -6.00460e-05, 0.00111925, -1.91348e-05, 2.97470e-05, 5.64197e-06, -4.86790e-06, -7.47441e-05, 5.01671e-06, -9.73478e-06, 4.71226e-05, 4.14501e-06, -6.19314e-05, -0.000990908, -1.55594e-05, 3.17897e-06, -5.64466e-06, 2.32544e-05, -1.92580e-05, -9.34792e-06, 0.000585389, -2.66625e-05, -1.51099e-05, 1.55506e-06, 2.92118e-05, -1.75075e-05, -5.53138e-05, 4.73336e-06, -0.000190178, -0.000407511, 0.00101097, -2.94065e-05, 1.09492e-05, 1.48063e-05, 0.000498114, 1.97546e-05, 8.48472e-06, -4.79723e-06, 0.000429315, -4.51327e-06, 6.23019e-06, -6.57924e-06, 2.14542e-06, -0.000116261, 3.04644e-06, 5.50514e-07, 2.13532e-05, 1.70614e-06, -2.44311e-05, -0.000288588, -5.14263e-06, -4.12850e-06, 2.07116e-06, 8.52461e-06, -8.02412e-06, -7.25291e-06, 0.000149062, -6.67296e-06, -1.67216e-06, 2.15586e-06, 9.47900e-06, 0.000765924, -0.000667760, 0.000136034, 0.000577858, 0.000282976, -0.000247622, -7.26225e-05, 0.000204668, 0.000346179, 0.00166702, 0.000421589, -0.000900229, 0.0184341, 0.0228037, -0.0416961, -0.000513640, -0.000885666, 0.00112785, -0.0208040, -0.000212215, -0.000703863, 0.000115796, -0.00754324, -0.000150840, -0.000964173, 0.00118126, 2.26409e-05, 0.0141106, -0.000496938, 0.000109460, 0.000192733, -1.54618e-05, -0.000487668, 0.00316075, -0.000397677, 4.85638e-05, -0.000247306, -0.000747009, 0.00108636, 0.000379286, -0.0158858, -0.000183161, 0.000538771, -0.000302495, -0.000464679, -0.00960941, -0.0292680, 0.0246750, -0.251069, 0.0330521, 0.00572265, -0.00691699, 0.0105749, -0.000753156, -0.00953489, -0.000832638, 0.0135597, 0.0504302, 0.00554654, 0.000653495, -0.000675914, -0.000802532, -0.00396353, 0.15478)
```


```
[info] fitMap qof = LinkedHashMap(rSq -> 0.963405, rSqBar -> 0.962474, sst -> 150761453.620609, sse -> 5517160.084849, mse0 -> 589.629164, rmse -> 24.282281, mae -> 15.939130, dfm -> 232.000000, df -> 9124.000000, fStat -> 1035.334785, aic -> -42657.469613, bic -> -40992.945572, maape -> 14.316041, smaape -> 14.720222)
```

```
[info]
```

```
[info] Run + title
```

SymRidgeRegressionXX: y actual, predicted

Reset Plot



R2: 0.963405


Adj R2: 0.962474

## Symbolic Lasso Regression:

```
Command Prompt - sbt
-07, -8.77799e-07, -4.46690e-05, 1.05184e-06, 3.30803e-07, -3.24808e-07, -2.16698e-07, -5.46940e-06, 6.
.63980e-07, 2.30039e-08, -6.22065e-07, 5.83632e-07, 4.88396e-07, 6.45258e-05, -2.23250e-06, -4.11175
e-06, -9.93039e-07, 5.01678e-07, -2.19406e-07, -9.41814e-06, 3.56195e-08, 5.81606e-07, 6.79823e-07, -
7.20755e-06, 5.82372e-08, 4.56029e-07, 2.27780e-07, 1.64346e-07, -8.43369e-06, 6.25641e-08, 1.67635e
-07, 3.54090e-07, -6.34009e-07, -2.18273e-07, 1.82878e-05, 3.05591e-08, -5.14233e-07, -1.35755e-08, -
7.19943e-05, -0.000139347, 1.74097e-05, 0.000226635, -0.000566950, 0.00221320, -6.87668e-05, 1.96829e
-05, 4.63924e-05, 0.00140633, 5.98926e-05, 2.32231e-05, -6.01596e-05, 0.00111664, -1.90033e-05, 2.
97596e-05, 5.75446e-06, -4.87193e-06, -9.97852e-05, 5.47455e-06, -9.77991e-06, 4.72374e-05, 4.07326e
-06, -6.22309e-05, -0.000993276, -1.54053e-05, 3.14698e-06, -5.71838e-06, 2.32148e-05, -1.91948e-05, -
0.31932e-06, 0.000591679, -2.66479e-05, -1.52117e-05, 1.53781e-06, 2.91178e-05, -1.74777e-05, -5.53214
e-05, 4.89859e-06, -0.000187005, -0.000412601, 0.00101726, -2.94200e-05, 1.09948e-05, 1.47590e-05, 0.
000503741, 1.97540e-05, 8.51693e-06, -4.88290e-06, 0.000429105, -4.51304e-06, 6.24593e-06, -6.50875
e-06, 2.15407e-06, -0.000123619, 3.19027e-06, 5.38604e-07, 2.13523e-05, 1.67259e-06, -2.45383e-05, -
0.000291689, -5.06880e-06, -4.12896e-06, 2.02488e-06, 8.50945e-06, -7.98094e-06, -7.26080e-06, 0.000151
088, -6.66459e-06, -1.70537e-06, 2.14015e-06, 9.45243e-06, 0.000768489, -0.000666137, 0.000142847, 0.
000607473, 0.000279200, -0.000247359, -7.17844e-05, 0.000201927, 0.000347427, 0.00166534, 0.000419
087, -0.000908427, 0.0183996, 0.0235226, -0.0420502, -0.000514440, -0.000901235, 0.00113020, -
0.0212816, -0.000212232, -0.000705894, 0.000116226, -0.00755183, -0.000150652, -0.000965714, 0.001178
80, 1.67862e-05, 0.0148371, -0.000509955, 0.000109538, 0.000192757, -1.28924e-05, -0.000485353, 0.
0317467, -0.000401686, 4.90522e-05, -0.000243230, -0.000747462, 0.00108668, 0.000378544, -0.01597
27, -0.000184830, 0.000540278, -0.000303063, -0.000463209, -0.00957925, -0.0292176, 0.0249894, -
0.249361, 0.0331281, 0.00574261, -0.00698223, 0.0105074, -0.000743263, -0.00954361, -0.00083
4453, 0.0135724, 0.0486696, 0.00561981, 0.000644675, -0.000694131, -0.000791242, -0.00398140, 0.
15309)
[info] fitMap qof = LinkedHashMap(rSq -> 0.963405, rSqBar -> 0.962478, sst -> 150761453.620609, sse -> 5517138.2
70910, mse0 -> 589.626832, rmse -> 24.282233, mae -> 15.939459, dfm -> 231.000000, df -> 9125.000000, fStat -> 1039.9349
87, aic -> -42659.451115, bic -> -41002.070954, mape -> 14.316473, smape -> 14.720733)
[info]
[info] Run + title
```

SymLassoRegressionXX: y actual, predicted

Reset Plot



**R2: 0.963405**

**Adj R2: 0.962478**

**Summary:** Both Symbolic Ridge Regression and Symbolic Lasso Regression are ideal for this model.

### **ForestFires Dataset:**

We have taken this dataset from the UCI repository. Our aim is to predict the burned area of the forest. There are 518 instances in the dataset.

1. X - x-axis spatial coordinate within the Montesinho park map: 1 to 9
2. Y - y-axis spatial coordinate within the Montesinho park map: 2 to 9
3. month - month of the year: 'jan' to 'dec'
4. day - day of the week: 'mon' to 'sun'
5. FPMC - FPMC index from the FWI system: 18.7 to 96.20
6. DMC - DMC index from the FWI system: 1.1 to 291.3
7. DC - DC index from the FWI system: 7.9 to 860.6
8. ISI - ISI index from the FWI system: 0.0 to 56.10
9. temp - temperature in Celsius degrees: 2.2 to 33.30
10. RH - relative humidity in %: 15.0 to 100
11. wind - wind speed in km/h: 0.40 to 9.40
12. rain - outside rain in mm/m2 : 0.0 to 6.4
13. area - the burned area of the forest (in ha): 0.00 to 1090.84  
(this output variable is very skewed towards 0.0, thus it may make sense to model with the logarithm transform).

### **Linear Regression:**

```
modelName mn = Linear Regression
```

```
[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac_QR, Fac_QR)))
```

```
[info] features fn = Array(intercept, region, accommodation_type, yearly_availability,  
minimum_nights, number_of_reviews, reviews_per_month, owned_hotels)
```


```
[info] parameter b = VectorD(291.366, -72.0484, -2.09576, 116.166,  
0.122837, -0.111580, -22.5625, -0.49375)
```

```
[info] fitMap qof = LinkedHashMap(rSq -> 0.048879, rSqBar -> 0.046552, sst ->  
473344895.742509, sse -> 450208477.787416, mse0 -> 156867.065431, rmse ->  
396.064471, mae -> 138.951817, dfm -> 7.000000, df -> 2862.000000, fStat -> 21.011354,  
aic -> -21223.479593, bic -> -21175.783054, mape -> 107.846900, smape -> 63.717290)
```

```
[info] fname = Array(intercept, region, accommodation_type, yearly_availability, minimum_nights, nu
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
[info] -----
[info] x0      291.365759      27.477024      10.603978      0.000000      NA
[info] x1      -72.048363      9.269240      -7.772845      0.000000      1.058020
[info] x2      -2.095756      10.805352      -0.193955      0.846211      1.187625
[info] x3      116.166033      17.014840      6.827336      0.000000      1.320475
[info] x4       0.122837      0.199793      0.614823      0.538672      1.049737
[info] x5      -0.111580      0.281115      -0.396920      0.691427      1.520657
[info] x6      -22.562545      7.349080      -3.070118      0.002140      1.620195
[info] x7      -0.493754      0.285724      -1.728079      0.083974      1.093946
[info] Residual standard error: 396.617634 on 2862.0 degrees of freedom
[info] Multiple R-squared:  0.048879,    Adjusted R-squared:  0.046552
[info] F-statistic: 21.011354089903588 on 7.0 and 2862.0 DF,  p-value: 0.0
[info] -----
[info]
```

## Cross-Validation

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | 0.037 | 0.061 | 0.048 | 0.010 | 0.013 |
[info] | rSqBar | 5 | 0.035 | 0.059 | 0.045 | 0.010 | 0.013 |
[info] | sst | 5 | 65472382.329 | 156684138.578 | 94645824.164 | 39425546.436 | 48962871.126 |
[info] | sse | 5 | 61458971.265 | 150647766.310 | 90394318.546 | 38469064.711 | 47775009.558 |
[info] | mse0 | 5 | 107071.379 | 262452.555 | 157481.391 | 67019.276 | 83231.724 |
[info] | rmse | 5 | 327.218 | 512.301 | 390.295 | 80.246 | 99.658 |
[info] | mae | 5 | 135.323 | 150.927 | 139.585 | 6.492 | 8.062 |
[info] | dfm | 5 | 7.000 | 7.000 | 7.000 | 0.000 | 0.000 |
[info] | df | 5 | 2862.000 | 2862.000 | 2862.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 15.819 | 26.699 | 20.504 | 4.732 | 5.877 |
[info] | aic | 5 | -4425.072 | -4140.791 | -4233.020 | 122.617 | 152.279 |
[info] | bic | 5 | -4390.251 | -4105.970 | -4198.199 | 122.617 | 152.279 |
[info] | mape | 5 | 101.210 | 115.381 | 108.296 | 5.461 | 6.782 |
[info] | smape | 5 | 61.940 | 66.447 | 64.042 | 1.729 | 2.148 |
[info] -----
```



## Lasso Regression:

[info] REPORT

```
[info] modelName mn = LassoRegression
```

```
[info] hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01)))
```

```
[info] features fn = Array(x0, x1, x2, x3, x4, x5, x6, x7, x8, x9, x10, x11)
```

```
[info] parameter b = VectorD(1.91525, 0.228364, 2.82000, 1.33261,
-0.268104, 0.101702, -0.0307685, -0.654885, 0.891355, -0.206911, 1.14547,
-3.1241)
```

```
[info] fitMap    qof = LinkedHashMap(rSq -> 0.025264, rSqBar -> 0.004032, sst ->
2090864.624009, sse -> 2038040.487307, mse0 -> 3942.051233, rmse -> 62.785757, mae ->
19.527744, dfm -> 11.000000, df -> 505.000000, fStat -> 1.189921, aic -> -2849.830723, bic ->
-2798.854208, mape -> Infinity, smape -> 153.844345)
```

[info] SUMMARY

```
[info] Parameters/Coefficients:
```

[info]	Var	Estimate	Std. Error	t value	Pr(> t )	VIF
--------	-----	----------	------------	---------	----------	-----

```
[info] -----
[info]      x0      1.015253      1.447138      1.323476      0.185677      NA
```

[info]	x0	1.915253	1.447138	1.323476	0.185677	NA
[info]	x1	0.228364	2.716092	0.084078	0.932994	1.426780

[info]	x1	0.228984	2.718892	0.004070	0.992994	1.428780
[info]	x2	2.819998	2.775993	1.015852	0.309700	5.103943

```
[info]      x3      1.332613      1.333233      0.999535      0.317535      1.014415
```

[info]	x4	-0.268104	0.275146	-0.974408	0.329854	0.294951
[info]	x5	0.101702	0.068030	1.404770	0.134074	2.427000

[info]	x5	0.101702	0.068039	1.494770	0.134974	2.427904
[info]	x6	-0.030769	0.031967	-0.962505	0.335796	8.040274

[info]	x6	-0.653765	0.651567	-0.962565	0.553758	0.648274
[info]	x7	-0.654885	0.703758	-0.930555	0.352084	1.316450

```
[info]      x8      0.891355      0.760351      1.172295      0.241079      2.492321
```

[info]	x9	-0.206911	0.207833	-0.995563	0.319462	1.470498
[info]	x10	1.145472	1.675054	0.683842	0.404075	1.151578

[info]	x10	1.145472	1.675054	0.683842	0.494075	1.151578
[info]	x11	-3.124122	9.614787	-0.324929	0.745235	1.035312

```
[info] Residual standard error: 63.527346 on 505.0 degrees of freedom
```

```
[info] Multiple R-squared:  0.025264,    Adjusted R-squared:  0.004032
```

```
[info] F-statistic: 1.1899214511000655 on 11.0 and 505.0 DF, p-value: 0.0
```

```
[info] -----
```

### Cross-validation:

```
[info] -----
```

```
[info] | showQofStatTable: Statistical Table for QoF |
```

```
[info] -----
[info] |      name |    num |      min |      max |      mean |      stdev |      interval |
[info] |-----|-----|-----|-----|-----|-----|-----|
```

```
[info] |      name |      num |      min |      max |      mean |      stdev |      interval |
[info] -----
```

[info]															
[info]		rSq		5		-1.223		0.022		-0.271		0.534		0.663	

[info]	rSqBar	5	-1.271	0.001	-0.299	0.546	0.678
--------	--------	---	--------	-------	--------	-------	-------

[info]	sst	5	22086.277	1319497.171	412859.609	559096.891	694346.471
[info]		5	18000.000	1220122.127	421127.800	560221.257	695712.020

[info]	sse	5	49089.909	1339122.427	421127.800	560221.257	695742.829
[info]	mse0	5	476.601	13001.189	4088.619	5439.041	6754.785

[info]	mse	5	476.601	15001.189	4688.619	3439.041	6734.783
[info]	rmse	5	21.831	114.023	52.451	40.888	50.779

[info]	mae	5	15.959	28.604	20.336	5.465	6.787
--------	-----	---	--------	--------	--------	-------	-------

[info]	dfm	5	11.000	11.000	11.000	0.000	0.000
--------	-----	---	--------	--------	--------	-------	-------

[info]	df	5	505.000	505.000	505.000	0.000	0.000
[info]	ss_tot	5	25.254	1.040	6.170	10.861	13.480

[info]	fStat	5	-25.254	1.049	-6.170	10.861	13.489
[info]	aic	5	-666.894	-503.269	-550.457	71.057	88.246

[info]	aic	5	-668.894	-565.269	-556.457	71.057	88.246
[info]	bic	5	-635.277	-471.652	-518.841	71.057	88.246


```
[info] |      mape |    5 | 1797693134862315700000000000000000000000000000000000000000000000000000
```

[illegible]

```
[info] |      smape |      5 |      145.773 |      164.441 |      156.293 |      6.881 |      8.546 |
[info] -----
```


```
[info] -----
```

LassoRegression: y actual, predicted  
Reset Plot



R^2 vs n for LassoRegression

— □ ×



## Ridge Regression:

[info] REPORT

[info] modelName mn = RidgeRegression

[info] hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01), factorization -> (Fac\_Cholesky,Fac\_Cholesky)))

[info] features fn = Array(X, Y, month, day, FPMC, DMC, DC, ISI, temp, RH, wind, rain)

[info] parameter b = VectorD(1.90078, 0.343908, 2.88462, 1.36215, -0.109384, 0.0967938, -0.0314179, -0.737874, 0.953676, -0.174952, 1.22324, -3.4021)

[info] fitMap qof = LinkedHashMap(rSq -> 0.025398, rSqBar -> 0.002193, sst ->

**2090864.624009, sse -> 2037760.331537, mse0 -> 3941.509345, rmse -> 62.781441, mae ->**

**19.550098, dfm -> 12.000000, df -> 504.000000, fStat -> 1.094525, aic -> -2847.795186, bic ->**


**-2792.570629, mape -> 155.410179, smape -> 138.169308)**

```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
[info] -----
[info] x0      1.900779      1.449515      1.311321      0.189749      NA
[info] x1      0.343908      2.753794      0.124885      0.900615      1.463962
[info] x2      2.884616      2.789390      1.034139      0.301071      5.143826
[info] x3      1.362152      1.339174      1.017158      0.309078      1.021589
[info] x4      -0.109384      0.662902      -0.165007      0.868938      1.708917
[info] x5      0.096794      0.070608      1.370858      0.170419      2.609922
[info] x6      -0.031418      0.032092      -0.979004      0.327578      8.088132
[info] x7      -0.737874      0.771744      -0.956112      0.339016      1.580170
[info] x8      0.953676      0.797038      1.196525      0.231492      2.733585
[info] x9      -0.174952      0.240871      -0.726330      0.467637      1.971519
[info] x10     1.223244      1.702437      0.718525      0.472433      1.187344
[info] x11     -3.402107      9.681662      -0.351397      0.725291      1.047829
[info] Residual standard error: 63.585967 on 504.0 degrees of freedom
[info] Multiple R-squared: 0.025398, Adjusted R-squared: 0.002193
[info] F-statistic: 1.0945253223939044 on 12.0 and 504.0 DF, p-value: 0.0
[info] -----
```


## Cross-validation:

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | -1.153 | 0.021 | -0.252 | 0.506 | 0.628 |
[info] | rSqBar | 5 | -1.204 | -0.003 | -0.282 | 0.518 | 0.643 |
[info] | sst | 5 | 22086.277 | 1319497.171 | 412859.609 | 559096.891 | 694346.471 |
[info] | sse | 5 | 47552.425 | 1327165.759 | 418344.845 | 555856.981 | 690322.803 |
[info] | mse0 | 5 | 461.674 | 12885.104 | 4061.600 | 5396.670 | 6702.163 |
[info] | rmse | 5 | 21.487 | 113.513 | 52.243 | 40.808 | 50.680 |
[info] | mae | 5 | 15.509 | 29.866 | 20.342 | 6.169 | 7.662 |
[info] | dfm | 5 | 12.000 | 12.000 | 12.000 | 0.000 | 0.000 |
[info] | df | 5 | 504.000 | 504.000 | 504.000 | 0.000 | 0.000 |
[info] | fStat | 5 | -22.493 | 0.892 | -5.312 | 9.767 | 12.130 |
[info] | aic | 5 | -663.393 | -501.068 | -548.105 | 70.513 | 87.571 |
[info] | bic | 5 | -629.142 | -466.816 | -513.853 | 70.513 | 87.571 |
[info] | mape | 5 | 113.445 | 296.321 | 175.665 | 70.923 | 88.080 |
[info] | smape | 5 | 126.783 | 156.154 | 140.360 | 12.063 | 14.981 |
[info] -----
```

RidgeRegression: y actual, predicted  
reset Plot




R<sup>2</sup> vs n for RidgeRegression with Forward






R<sup>2</sup> vs n for RidgeRegression with Backward



R<sup>2</sup> vs n for RidgeRegression with Stepwise



## Quadratic Regression:

[info] REPORT

[info] modelName mn = PolyRegression

[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac\_Cholesky,Fac\_Cholesky)))

[info] features fn = Array(x0, x1, x2)

[info] parameter b = VectorD(3.39145, 0.0459032, 0.025567)

[info] fitMap qof = LinkedHashMap(rSq -> **0.301324**, rSqBar -> **0.298605**, sst ->

**780.529981**, sse -> **545.337524**, mse0 -> **1.054811**, rmse -> **1.027040**, mae -> **0.823228**, dfm


-> 2.000000, df -> 514.000000, fStat -> 110.838625, aic -> -741.385309, bic -> -728.641180,

mape -> 22.496013, smape -> 20.272131)

```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate   Std. Error    t value      Pr(>|t|)      VIF
[info] -----
[info] x0      3.391446      0.187421     18.095352    0.000000      NA
[info] x1      0.045903      0.090048     0.509766    0.610215     21.112255
[info] x2      0.025567      0.009328     2.740913    0.006127     21.112255
[info] Residual standard error: 1.030033 on 514.0 degrees of freedom
[info] Multiple R-squared: 0.301324, Adjusted R-squared: 0.298605
[info] F-statistic: 110.83862513206601 on 2.0 and 514.0 DF, p-value: 0.0
[info] -----
```

## Cross-Validation:

```
[info] |-----|
[info] | showQofStatTable: Statistical Table for QoF |
[info] |-----|
[info] | name | num | min | max | mean | stdev | interval |
[info] |-----|
[info] | rSq | 5 | 0.160 | 0.370 | 0.283 | 0.079 | 0.098 |
[info] | rSqBar | 5 | 0.157 | 0.367 | 0.280 | 0.079 | 0.099 |
[info] | sst | 5 | 116.718 | 188.718 | 154.548 | 26.867 | 33.366 |
[info] | sse | 5 | 94.843 | 127.205 | 109.723 | 15.255 | 18.945 |
[info] | mse0 | 5 | 0.921 | 1.235 | 1.065 | 0.148 | 0.184 |
[info] | rmse | 5 | 0.960 | 1.111 | 1.030 | 0.071 | 0.089 |
[info] | mae | 5 | 0.757 | 0.916 | 0.826 | 0.069 | 0.086 |
[info] | dfm | 5 | 2.000 | 2.000 | 2.000 | 0.000 | 0.000 |
[info] | df | 5 | 514.000 | 514.000 | 514.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 49.037 | 150.724 | 104.832 | 38.029 | 47.229 |
[info] | aic | 5 | -151.696 | -136.356 | -143.409 | 7.231 | 8.980 |
[info] | bic | 5 | -143.792 | -128.452 | -135.505 | 7.231 | 8.980 |
[info] | mape | 5 | 19.666 | 26.336 | 22.566 | 2.446 | 3.038 |
[info] | smape | 5 | 18.127 | 23.160 | 20.341 | 1.978 | 2.457 |
[info] |-----|
```



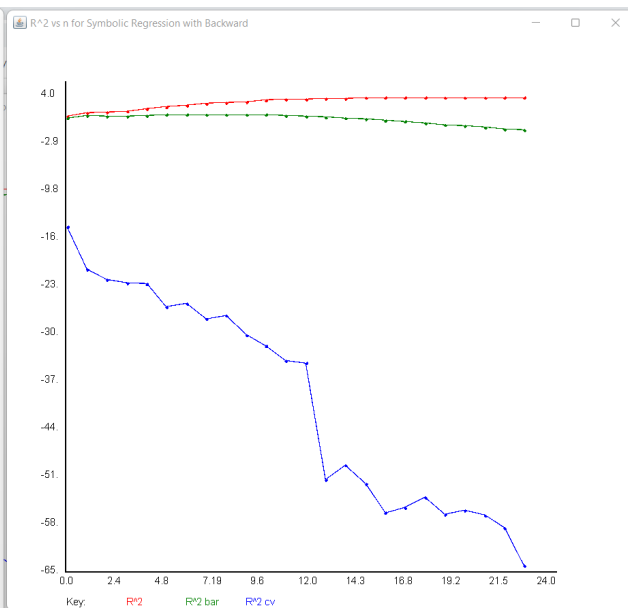


[info] REPORT


```
[info] modelName mn = SymbolicRegression.quadratic
[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac_QR,Fac_QR)))
[info] features fn = Array(intercept, X, Y, month, day, FPMC, DMC, DC, ISI, temp, RH, wind,
rain, X^2, Y^2, month^2, day^2, FPMC^2, DMC^2, DC^2, ISI^2, temp^2, RH^2, wind^2, rain^2)
[info] parameter b = VectorD(33.8738, -0.252232, -1.34350, -13.8009, -4.86745,
-0.660461, 0.531026, -0.0637533, -1.45247, 0.612062, -0.461598, 8.35282,
-23.8968, 0.282141, 0.110506, 1.27133, 0.814943, 0.00509256,
-0.00131266, 2.27165e-05, 0.0127240, 0.0142198, 0.00354986, -0.757032,
3.3328)
[info] fitMap qof = LinkedHashMap(rSq -> 0.036853, rSqBar -> -0.010130, sst ->
2090864.624009, sse -> 2013810.569610, mse0 -> 3895.184854, rmse -> 62.411416, mae ->
19.712527, dfm -> 24.000000, df -> 492.000000, fStat -> 0.784388, aic -> -2820.739045, bic ->
-2714.537973, mape -> Infinity, smape -> 154.562640)
```

[info]	SUMMARY					
[info]	Parameters/Coefficients:					
[info]	Var	Estimate	Std. Error	t value	Pr(> t )	VIF
[info]	x0	33.873753	121.338933	0.279166	0.780117	NA
[info]	x1	-0.252232	6.614921	-0.038131	0.969583	29.531729
[info]	x2	-1.343496	10.977560	-0.122386	0.902594	32.979879
[info]	x3	-13.080879	14.987043	-0.920854	0.357127	146.679318
[info]	x4	-4.867455	6.501757	-0.748637	0.454076	23.786627
[info]	x5	-0.660461	3.495172	-0.188964	0.850121	46.927668
[info]	x6	0.531026	0.310825	1.708442	0.087554	49.959470
[info]	x7	-0.063753	0.123401	-0.516634	0.605412	118.132937
[info]	x8	-1.452467	1.918072	-0.757253	0.448898	9.641757
[info]	x9	0.612062	3.236598	0.189107	0.850009	44.526772
[info]	x10	-0.461598	0.990324	-0.466108	0.641138	32.919718
[info]	x11	8.352821	6.514945	1.282102	0.199807	17.176129
[info]	x12	-23.896795	38.813508	-0.615682	0.538104	16.635101
[info]	x13	0.282141	0.697240	0.404654	0.685732	30.57021
[info]	x14	0.110506	1.187648	0.093046	0.925867	23.635443
[info]	x15	1.271333	1.054702	1.205395	0.228051	112.273831
[info]	x16	0.814943	0.806983	1.009863	0.312561	23.589064
[info]	x17	0.005093	0.025958	0.196184	0.844466	56.521938
[info]	x18	-0.001313	0.000931	-1.409855	0.158582	33.007449
[info]	x19	0.000023	0.000108	0.210328	0.833421	64.446159
[info]	x20	0.012724	0.041757	0.304714	0.760584	5.590794
[info]	x21	0.014220	0.075739	0.187748	0.851074	33.012194
[info]	x22	0.003550	0.000335	0.380286	0.703733	32.007929
[info]	x23	-0.757032	0.709920	-1.066362	0.286260	17.842384
[info]	x24	3.332823	6.268611	0.531668	0.594956	16.123072
[info]	Residual standard error: 63.977425 on 492.0 degrees of freedom					
[info]	Multiple R-squared: 0.036853, Adjusted R-squared: -0.010130					
[info]	F-statistic: 0.7843876375582628 on 24.0 and 492.0 DF, p-value: 0.0					

### Cross-validation:

[illegible]





## Symbolic Lasso Regression:

[info] REPORT

[info] modelName mn = SymLassoRegression.quadratic

[info] hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01)))



[info] features fn = Array(X, Y, month, day, FFMC, DMC, DC, ISI, temp, RH, wind, rain, X^2, Y^2, month^2, day^2, FFMC^2, DMC^2, DC^2, ISI^2, temp^2, RH^2, wind^2, rain^2)


[info] parameter b = VectorD(-0.0338159, -0.800841, -14.5594, -4.67453, 0.199852, 0.527027, -0.0573202, -1.36578, 0.822753, -0.487909, 8.34073, -24.5477, 0.255144, 0.0611386, 1.31842, 0.792161, -0.000661795, -0.00129973, 1.69202e-05, 0.0115908, 0.0105412, 0.00400019, -0.757373, 3.4301)

[info] fitMap qof = LinkedHashMap(rSq -> 0.036700, rSqBar -> -0.008241, sst -> 2090864.624009, sse -> 2014129.561691, mse0 -> 3895.801860, rmse -> 62.416359, mae -> 19.750481, dfm -> 23.000000, df -> 493.000000, fStat -> 0.816630, aic -> -2822.779989, bic -> -2720.826960, mape -> Infinity, smape -> 154.122194)

```
[info] fname = Array(X, Y, month, day, FFMC, DMC, DC, ISI, temp, RH, wind, rain, X^2, Y^2, month^2, day^2, FFMC^2, DMC^2, DC^2, ISI^2, temp^2, RH^2, wind^2, rain^2)
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var Estimate Std. Error t value Pr(>|t|) VIF
-----
[info] x0 -0.033816 6.562345 -0.005153 0.995889 NA
[info] x1 -0.800841 10.794149 -0.074192 0.940857 22.260038
[info] x2 -14.559415 14.724352 -0.988798 0.322762 141.847722
[info] x3 -4.674529 6.458901 -0.723735 0.469229 23.518071
[info] x4 0.199852 1.647390 0.121315 0.903442 10.444738
[info] x5 0.527027 0.310205 1.698964 0.089326 49.853750
[info] x6 -0.057320 0.121116 -0.473265 0.636024 114.011829
[info] x7 -1.365778 1.891004 -0.722259 0.470141 9.389197
[info] x8 0.822753 3.144325 0.261663 0.793581 42.102879
[info] x9 -0.487909 0.984904 -0.495387 0.620327 32.621435
[info] x10 8.340726 6.508709 1.281472 0.200028 17.175386
[info] x11 -24.547700 38.706231 -0.634205 0.525947 16.574271
[info] x12 0.255144 0.689856 0.369851 0.711493 29.986954
[info] x13 0.061139 1.173324 0.052107 0.958443 23.111970
[info] x14 1.318423 1.040119 1.267569 0.204952 109.395071
[info] x15 0.792161 0.802098 0.987611 0.323343 23.347995
[info] x16 -0.000662 0.015763 -0.041984 0.966512 20.881669
[info] x17 -0.001300 0.000929 -1.399001 0.161813 32.925981
[info] x18 0.000017 0.000106 0.159700 0.873047 62.063511
[info] x19 0.011591 0.041521 0.279157 0.780124 5.537981
[info] x20 0.010541 0.074512 0.141469 0.887499 32.011619
[info] x21 0.004000 0.009186 0.435487 0.663209 31.138735
[info] x22 -0.757373 0.709255 -1.067844 0.285591 17.842328
[info] x23 3.430108 6.252925 0.548561 0.583307 16.072544
[info] Residual standard error: 63.917568 on 493.0 degrees of freedom
[info] Multiple R-squared: 0.036700, Adjusted R-squared: -0.008241
[info] F-statistic: 0.8166303749967986 on 23.0 and 493.0 DF, p-value: 0.0
```

### Cross-validation:

[illegible]



## Symbolic Ridge Regression:

[info] REPORT

[info] modelName mn = SymRidgeRegression.quadratic

[info] hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01), factorization -> (Fac\_Cholesky,Fac\_Cholesky)))

[info] features fn = Array(X, Y, month, day, FPMC, DMC, DC, ISI, temp, RH, wind, rain, X^2, Y^2, month^2, day^2, FPMC^2, DMC^2, DC^2, ISI^2, temp^2, RH^2, wind^2, rain^2)

[info] parameter b = VectorD(-0.0344328, -0.800701, -14.5488, -4.67538, 0.200168, 0.526929, -0.0573334, -1.36557, 0.821176, -0.487953, 8.34000, -24.4586, 0.255175, 0.0611479, 1.31768, 0.792268, -0.000664959, -0.00129945, 1.69325e-05, 0.0115876, 0.0105687, 0.00399885, -0.757376, 3.4162)

[info] fitMap qof = LinkedHashMap(rSq -> 0.036700, rSqBar -> -0.010290, sst -> 2090864.624009, sse -> 2014129.586689, mse0 -> 3895.801908, rmse -> 62.416359, mae -> 19.749900, dfm -> 24.000000, df -> 492.000000, fStat -> 0.781016, aic -> -2820.779992, bic -> -2714.578920, mape -> Infinity, smape -> 154.123007)



```
[info] Run + title
[info] fname = Array(X, Y, month, day, FPMC, DMC, DC, ISI, temp, RH, wind, rain, X^2, Y^2, month^2, day^2, FPMC^2, DMC^2, DC^2, ISI^2, temp^2, RH^2, wind^2, rain^2)
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate  Std. Error  t value    Pr(>|t|)    VIF
[info] -----
[info] x0      -0.034433   6.569011   -0.005242   0.995818    NA
[info] x1      -0.800701   10.805113   -0.074104   0.940928    22.260038
[info] x2     -14.548803   14.739308   -0.987075   0.323606    141.847722
[info] x3      -4.675380    6.454541   -0.723132   0.469599    23.510071
[info] x4      0.200168    1.649063    0.121383   0.903388    10.444738
[info] x5      0.526929    0.310520    1.696922   0.089711    49.853750
[info] x6     -0.057333    0.121239   -0.472894   0.636289    114.011829
[info] x7     -1.365569    1.802925   -0.721407   0.470659    9.389167
[info] x8      0.821176    3.147519    0.260896   0.794173    42.102879
[info] x9     -0.487953    0.985905   -0.494929   0.620650    32.621435
[info] x10     8.340004    6.515320    1.280061    0.208524    17.175386
[info] x11     -24.458638   39.745547   -0.613163   0.537868    16.574771
[info] x12     0.255175    0.690557    0.369520    0.711740    29.986954
[info] x13     0.061148    1.174515    0.052062   0.958479    23.111970
[info] x14     1.317678    1.041175    1.265666   0.205568    109.395071
[info] x15     0.792268    0.802513    0.985742    0.323769    23.547995
[info] x16     -0.000665    0.015779   -0.042142   0.966386    20.881669
[info] x17     -0.001299    0.000930   -1.397279   0.162330    32.925981
[info] x18     0.000017    0.000106    0.159744    0.873003    62.063511
[info] x19     0.011588    0.041563    0.278799    0.780399    5.537981
[info] x20     0.010569    0.074588    0.141695    0.887321    32.011619
[info] x21     0.003999    0.009195    0.434899    0.663636    31.138735
[info] x22     -0.757376    0.780975   -0.969655   0.336078    17.042328
[info] x23     3.416219    6.259277    0.545785    0.585214    16.072544
[info] Residual standard error: 63.982492 on 492.0 degrees of freedom
[info] Multiple R-squared: 0.036700, Adjusted R-squared: -0.010290
[info] F-statistic: 0.781016 on 24.0 and 492.0 Df, p-value: 0.0
[info]
```

### Cross-Validation:


[illegible]

 SymRidgeRegression.quadratic: y actual, predicted

Reset Plot







## Expedia Dataset:

### Regression:

[info] REPORT

[info] modelName mn = Regression

[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac\_QR,Fac\_QR)))

[info] features fn = Array(intercept, region, accommodation\_type, yearly\_availability, minimum\_nights, number\_of\_reviews, reviews\_per\_month, owned\_hotels)

[info] parameter b = VectorD(291.366, -72.0484, -2.09576, 116.166, 0.122837, -0.111580, -22.5625, -0.49375)

[info] fitMap qof = LinkedHashMap(rSq -> **0.048879**, rSqBar -> **0.046552**, sst -> **473344895.742509**, sse -> **450208477.787416**, mse0 -> **156867.065431**, rmse -> **396.064471**, mae -> **138.951817**, dfm -> 7.000000, df -> 2862.000000, fStat -> 21.011354, aic -> -21223.479593, bic -> -21175.783054, mape -> 107.846900, smape -> 63.717290)

[info] SUMMARY

[info] Parameters/Coefficients:

Var	Estimate	Std. Error	t value	Pr(> t )	VIF
x0	291.365759	27.477024	10.603978	0.000000	NA
x1	-72.048363	9.269240	-7.772845	0.000000	1.058020
x2	-2.095756	10.805352	-0.193955	0.846211	1.187625
x3	116.166033	17.014840	6.827336	0.000000	1.320475
x4	0.122837	0.199793	0.614823	0.538672	1.049737
x5	-0.111580	0.281115	-0.396920	0.691427	1.520657
x6	-22.562545	7.349080	-3.070118	0.002140	1.620195
x7	-0.493754	0.285724	-1.728079	0.083974	1.093946

[info] Residual standard error: 396.617634 on 2862.0 degrees of freedom

[info] Multiple R-squared: 0.048879, Adjusted R-squared: 0.046552


[info] F-statistic: 21.011354089903588 on 7.0 and 2862.0 DF, p-value: 0.0

[info] -----

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] |   name | num |   min |   max |   mean |   stdev | interval |
[info] -----
[info] |   rSq |   5 |   0.037 |   0.061 |   0.048 |   0.010 |   0.013 |
[info] | rSqBar |   5 |   0.035 |   0.059 |   0.045 |   0.010 |   0.013 |
[info] |   sst |   5 | 65472382.329 | 156684138.578 | 94645824.164 | 39425546.436 | 48962871.126 |
[info] |   sse |   5 | 61458971.265 | 150647766.310 | 90394318.546 | 38469064.711 | 47775009.558 |
[info] |  mse0 |   5 | 107071.379 | 262452.555 | 157481.391 | 67019.276 | 83231.724 |
[info] |   rmse |   5 |   327.218 |   512.301 |   390.295 |    80.246 |   99.658 |
[info] |   mae |   5 |   135.323 |   150.927 |   139.585 |    6.492 |    8.062 |
[info] |   dfm |   5 |    7.000 |    7.000 |    7.000 |    0.000 |    0.000 |
[info] |   df |   5 | 2862.000 | 2862.000 | 2862.000 |    0.000 |    0.000 |
[info] |  fStat |   5 |   15.819 |   26.699 |   20.504 |    4.732 |    5.877 |
[info] |   aic |   5 | -4425.072 | -4140.791 | -4233.020 |   122.617 |   152.279 |
[info] |   bic |   5 | -4390.251 | -4105.970 | -4198.199 |   122.617 |   152.279 |
[info] |   mape |   5 |   101.210 |   115.381 |   108.296 |    5.461 |    6.782 |
[info] |   smape |   5 |    61.940 |    66.447 |    64.042 |    1.729 |    2.148 |
[info] -----
```

Regression: y actual, predicted

Reset Plot



Linear Regression:

Multiple R-squared: 0.048879

Adjusted R-squared: 0.46552

### Lasso Regression:

[info] REPORT

[info] modelName mn = LassoRegression

[info] hparameter hp = HyperParameter(HashMap(lambda -> (0.01,0.01)))

[info] features fn = Array(x0, x1, x2, x3, x4, x5, x6)

[info] parameter b = VectorD(-16.6993, 75.5694, 116.332, 0.349717, 0.290970, -6.20799, 0.090560)


[info] fitMap qof = LinkedHashMap(rSq -> 0.011510, rSqBar -> 0.009439, sst -> 473344895.742509, sse -> 467896598.283164, mse0 -> 163030.173618, rmse -> 403.769951,

mae -> **151.890328**, dfm -> 6.000000, df -> 2863.000000, fStat -> 5.556240, aic -> -21280.784992, bic -> -21239.050521, mape -> 117.542858, smape -> 75.067564)

```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
[info] -----
[info] x0      -16.699272      7.807558      -2.138860      0.032447      NA
[info] x1       75.569360      8.097478      9.332456      0.000000      0.641973
[info] x2      116.332194      17.342829      6.707798      0.000000      1.320474
[info] x3       0.349717      0.202473      1.727228      0.084127      1.037698
[info] x4       0.290970      0.283909      1.024871      0.305424      1.492926
[info] x5       -6.207991      7.323934      -0.847631      0.396644      1.548837
[info] x6       0.090560      0.285765      0.316905      0.751316      1.053259
[info] Residual standard error: 404.263256 on 2863.0 degrees of freedom
[info] Multiple R-squared:  0.011510,    Adjusted R-squared:  0.009439
[info] F-statistic: 5.55624030442443 on 6.0 and 2863.0 DF,  p-value: 0.0
[info] -----
```

## Cross-Validation:

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | -0.016 | 0.024 | 0.006 | 0.016 | 0.020 |
[info] | rSqBar | 5 | -0.018 | 0.022 | 0.004 | 0.016 | 0.020 |
[info] | sst | 5 | 65472382.329 | 156684138.578 | 94645824.164 | 39425546.436 | 48962871.126 |
[info] | sse | 5 | 63887089.643 | 153988041.626 | 93885029.824 | 38345129.817 | 47621093.917 |
[info] | mse0 | 5 | 111301.550 | 268271.850 | 163562.770 | 66803.362 | 82963.578 |
[info] | rmse | 5 | 333.619 | 517.950 | 398.269 | 78.615 | 97.633 |
[info] | mae | 5 | 146.119 | 156.987 | 152.344 | 4.411 | 5.479 |
[info] | dfm | 5 | 6.000 | 6.000 | 6.000 | 0.000 | 0.000 |
[info] | df | 5 | 2863.000 | 2863.000 | 2863.000 | 0.000 | 0.000 |
[info] | fStat | 5 | -7.317 | 11.840 | 2.938 | 7.812 | 9.702 |
[info] | aic | 5 | -4430.735 | -4153.643 | -4245.897 | 117.925 | 146.452 |
[info] | bic | 5 | -4400.267 | -4123.174 | -4215.429 | 117.925 | 146.452 |
[info] | mape | 5 | 110.991 | 127.404 | 117.919 | 7.633 | 9.480 |
[info] | smape | 5 | 70.770 | 79.015 | 75.273 | 3.337 | 4.145 |
[info] -----
```



## Lasso Regression:

Multiple R-squared: 0.011510  
Adjusted R-squared: 0.009439

## Ridge Regression:

[info] REPORT

[info] modelName mn = RidgeRegression

[info] hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01), factorization -> (Fac\_Cholesky,Fac\_Cholesky)))

[info] features fn = Array(region, accommodation\_type, yearly\_availability, minimum\_nights, number\_of\_reviews, reviews\_per\_month, owned\_hotels)


[info] parameter b = VectorD(-72.0478, -2.09529, 116.164, 0.122841, -0.111580, -22.5627, -0.49374)

[info] fitMap qof = LinkedHashMap(rSq -> **0.048879**, rSqBar -> **0.046552**, sst -> **473344895.742509**, sse -> **450208477.790368**, mse0 -> **156867.065432**, rmse -> **396.064471**, mae -> **138.951617**, dfm -> 7.000000, df -> 2862.000000, fStat -> 21.011354, aic -> -21223.479593, bic -> -21175.783054, mape -> 331.066601, smape -> 128.639245)

```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate  Std. Error    t value      Pr(>|t|)      VIF
[info] -----
[info] x0      -72.047806    9.269240     -7.772785    0.000000      NA
[info] x1       -2.095288    10.805352    -0.193912    0.846245      1.187625
[info] x2      116.163820    17.014840     6.827206    0.000000      1.320475
[info] x3       0.122841     0.199793     0.614843    0.538658      1.049737
[info] x4      -0.111580     0.281115    -0.396920    0.691427      1.520657
[info] x5      -22.562680     7.349080    -3.070137    0.002140      1.620195
[info] x6      -0.493745     0.285724    -1.728049    0.083980      1.093946
[info] Residual standard error: 396.617634 on 2862.0 degrees of freedom
[info] Multiple R-squared:  0.048879, Adjusted R-squared:  0.046552
[info] F-statistic: 21.01135408708497 on 7.0 and 2862.0 DF, p-value: 0.0
[info] -----
```

## Cross-Validation:

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | 0.037 | 0.061 | 0.048 | 0.010 | 0.013 |
[info] | rSqBar | 5 | 0.035 | 0.059 | 0.046 | 0.010 | 0.013 |
[info] | sst | 5 | 65472382.329 | 156684138.578 | 94645824.164 | 39425546.436 | 48962871.126 |
[info] | sse | 5 | 61458775.771 | 150606394.364 | 90371609.659 | 38459384.951 | 47762988.194 |
[info] | mse0 | 5 | 107071.038 | 262380.478 | 157441.829 | 67002.413 | 83210.781 |
[info] | rmse | 5 | 327.217 | 512.231 | 390.246 | 80.236 | 99.646 |
[info] | mae | 5 | 135.220 | 151.965 | 139.533 | 7.044 | 8.748 |
[info] | dfm | 5 | 7.000 | 7.000 | 7.000 | 0.000 | 0.000 |
[info] | df | 5 | 2862.000 | 2862.000 | 2862.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 15.891 | 26.701 | 20.613 | 4.717 | 5.858 |
[info] | aic | 5 | -4424.941 | -4140.790 | -4232.947 | 122.586 | 152.240 |
[info] | bic | 5 | -4390.120 | -4105.969 | -4198.126 | 122.586 | 152.240 |
[info] | mape | 5 | 233.764 | 420.446 | 331.829 | 71.912 | 89.308 |
[info] | smape | 5 | 122.942 | 132.594 | 128.794 | 4.127 | 5.126 |
[info] -----
```



Ridge Regression:

Multiple R-squared: 0.048879

Adjusted R-squared: 0.046552

**Symbolic Regression:**

[info] REPORT

[info] modelName mn = SymbolicRegression.quadratic

[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac\_QR,Fac\_QR)))

[info] features fn = Array(intercept, region, accommodation\_type, yearly\_availability, minimum\_nights, number\_of\_reviews, reviews\_per\_month, owned\_hotels, region^2, accommodation\_type^2, yearly\_availability^2, minimum\_nights^2, number\_of\_reviews^2, reviews\_per\_month^2, owned\_hotels^2)

[info] parameter b = VectorD(380.487, -170.932, 28.3262, -5.56516e+13, 0.350423, -0.250481, -60.1795, -2.05479, 22.5090, -7.89439, 5.56516e+13, -0.000540263, 0.00153794, 7.18678, 0.0065093)

[info] fitMap qof = LinkedHashMap(rSq -> **0.056665**, rSqBar -> **0.052039**, sst -> **473344895.742509**, sse -> **446522666.146301**, mse0 -> **155582.810504**, rmse -> **394.439869**, mae -> **139.477534**, dfm -> 14.000000, df -> 2855.000000, fStat -> 12.249813, aic -> -21197.683037, bic -> -21108.252027, mape -> 106.793957, smape -> 65.310552)


```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var Estimate Std. Error t value Pr(>|t|) VIF
[info] -----
[info] x0 380.486725 NaN NaN 0.000000 NA
[info] x1 -170.932022 NaN NaN 0.000000 17.309176
[info] x2 28.326164 NaN NaN 0.000000 79.760563
[info] x3 -55651632671836.130000 NaN NaN 0.000000 Infinity
[info] x4 0.350423 NaN NaN 0.000000 3.543752
[info] x5 -0.250481 NaN NaN 0.000000 6.196127
[info] x6 -60.179486 NaN NaN 0.000000 8.128142
[info] x7 -2.054790 NaN NaN 0.000000 5.334912
[info] x8 22.508976 NaN NaN 0.000000 17.134495
[info] x9 -7.894391 NaN NaN 0.000000 73.409289
[info] x10 55651632671945.820000 NaN NaN 0.000000 Infinity
[info] x11 -0.000540 NaN NaN 0.000000 3.379866
[info] x12 0.001538 NaN NaN 0.000000 4.664131
[info] x13 7.186780 NaN NaN 0.000000 6.059455
[info] x14 0.006509 NaN NaN 0.000000 4.898707
[info] Residual standard error: 395.474694 on 2855.0 degrees of freedom
[info] Multiple R-squared: 0.056665, Adjusted R-squared: 0.052039
[info] F-statistic: 12.24981256000395 on 14.0 and 2855.0 DF, p-value: 0.0
[info] -----
```

## Cross-validation:

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | 0.033 | 0.075 | 0.051 | 0.018 | 0.023 |
[info] | rSqBar | 5 | 0.028 | 0.070 | 0.047 | 0.018 | 0.023 |
[info] | sst | 5 | 65472382.329 | 156684138.578 | 94645824.164 | 39425546.436 | 48962871.126 |
[info] | sse | 5 | 60562171.612 | 149582876.103 | 90067557.337 | 38327769.315 | 47599533.784 |
[info] | mse0 | 5 | 105509.010 | 260597.345 | 156912.121 | 66773.117 | 82926.017 |
[info] | rmse | 5 | 324.822 | 510.487 | 389.555 | 80.307 | 99.734 |
[info] | mae | 5 | 135.179 | 150.843 | 141.401 | 5.940 | 7.377 |
[info] | dfm | 5 | 14.000 | 14.000 | 14.000 | 0.000 | 0.000 |
[info] | df | 5 | 2855.000 | 2855.000 | 2855.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 6.874 | 16.534 | 11.124 | 4.160 | 5.167 |
[info] | aic | 5 | -4409.254 | -4123.167 | -4217.989 | 123.175 | 152.972 |
[info] | bic | 5 | -4343.965 | -4057.877 | -4152.699 | 123.175 | 152.972 |
[info] | mape | 5 | 104.255 | 114.963 | 108.579 | 4.151 | 5.156 |
[info] | smape | 5 | 64.422 | 67.400 | 66.332 | 1.392 | 1.728 |
[info] -----
```

SymbolicRegression.quadratic: y actual, predicted

Reset Plot



## Symbolic Regression:

Multiple R-squared: 0.056665

Adjusted R-squared: 0.052039

## Quadratic Regression:

[info] REPORT

[info] modelName mn = PolyRegression

[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac\_Cholesky,Fac\_Cholesky)))

[info] features fn = Array(x0, x1, x2)

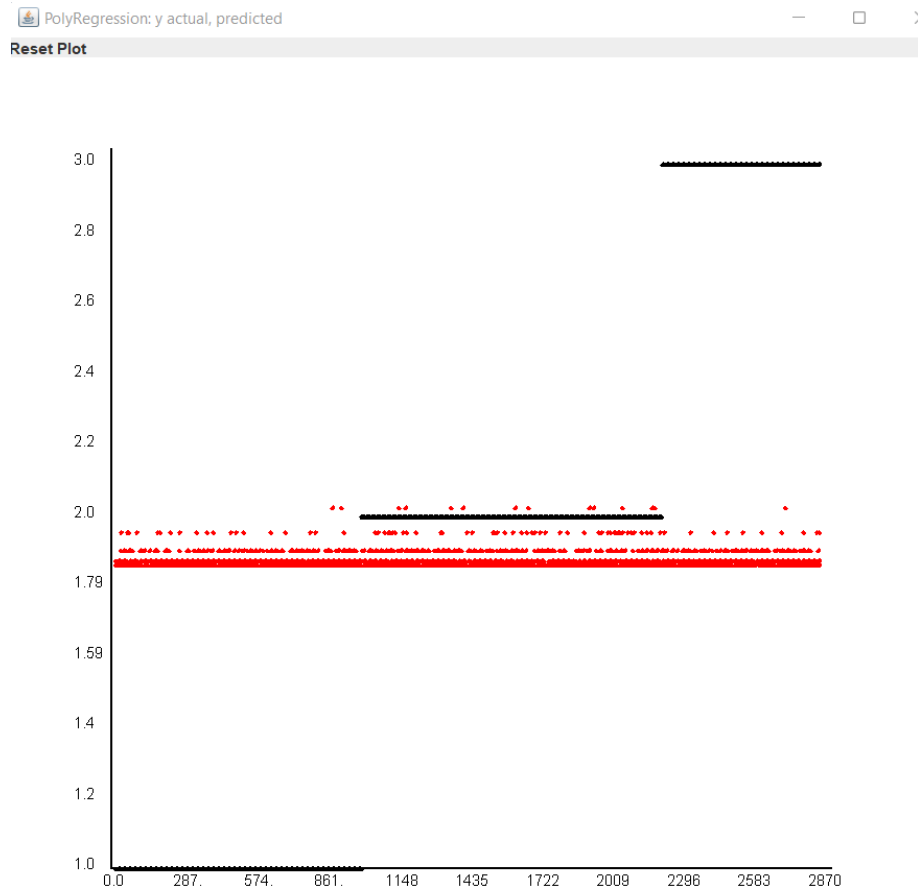
[info] parameter b = VectorD(1.87193, -0.0197617, 0.010120)

[info] fitMap qof = LinkedHashMap(rSq -> **0.000839**, rSqBar -> **0.000142**, sst -> **1600.093728**, sse -> **1598.751508**, mse0 -> **0.557056**, rmse -> **0.746362**, mae -> **0.609821**, dfm -> 2.000000, df -> 2867.000000, fStat -> 1.203485, aic -> -3226.750857, bic -> -3208.864655, mape -> 41.527495, smape -> 34.293227)

```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
[info] -----
[info] x0      1.871927      0.067571      27.703297      0.000000      NA
[info] x1     -0.019762      0.069351     -0.284951      0.775682     16.707235
[info] x2      0.010121      0.015585      0.649397      0.516082     16.707235
[info] Residual standard error: 0.746752 on 2867.0 degrees of freedom
[info] Multiple R-squared:  0.000839, Adjusted R-squared:  0.000142
[info] F-statistic: 1.2034849518511246 on 2.0 and 2867.0 DF, p-value: 0.0
[info] -----
[info]
```

## Cross-Validation:

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | -0.003 | 0.000 | -0.002 | 0.001 | 0.002 |
[info] | rSqBar | 5 | -0.004 | -0.000 | -0.003 | 0.001 | 0.002 |
[info] | sst | 5 | 313.401 | 325.814 | 319.778 | 4.820 | 5.985 |
[info] | sse | 5 | 314.133 | 326.737 | 320.371 | 4.755 | 5.905 |
[info] | mse0 | 5 | 0.547 | 0.569 | 0.558 | 0.008 | 0.010 |
[info] | rmse | 5 | 0.740 | 0.754 | 0.747 | 0.006 | 0.007 |
[info] | mae | 5 | 0.601 | 0.620 | 0.610 | 0.007 | 0.009 |
[info] | dfm | 5 | 2.000 | 2.000 | 2.000 | 0.000 | 0.000 |
[info] | df | 5 | 2867.000 | 2867.000 | 2867.000 | 0.000 | 0.000 |
[info] | fStat | 5 | -4.052 | 0.632 | -2.655 | 1.868 | 2.319 |
[info] | aic | 5 | -646.821 | -635.508 | -641.107 | 4.268 | 5.300 |
[info] | bic | 5 | -633.763 | -622.450 | -628.049 | 4.268 | 5.300 |
[info] | mape | 5 | 40.450 | 43.270 | 41.570 | 1.068 | 1.326 |
[info] | smape | 5 | 33.906 | 35.196 | 34.316 | 0.515 | 0.639 |
[info] -----
```



Ridge Regression:

Multiple R-squared: 0.000839

Adjusted R-squared: 0.000142

**Symbolic Lasso Regression:**

[info] REPORT

[info] modelName mn = SymLassoRegression.quadratic

[info] hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01)))


[info] features fn = Array(region, accommodation\_type, yearly\_availability, minimum\_nights, number\_of\_reviews, reviews\_per\_month, owned\_hotels, region^2, accommodation\_type^2, yearly\_availability^2, minimum\_nights^2, number\_of\_reviews^2, reviews\_per\_month^2, owned\_hotels^2)

[info] parameter b = VectorD(-93.5917, 376.349, 26.1912, 0.373307, -0.130269, -44.6193, -1.65233, 6.67263, -90.2034, 26.1912, -0.000582570, 0.000987059, 5.30805, 0.0054025)

[info] fitMap qof = LinkedHashMap(rSq -> **0.050100**, rSqBar -> **0.045776**, sst -> **473344895.742509**, sse -> **449630542.789972**, mse0 -> **156665.694352**, rmse -> **395.810175**, mae -> **139.209646**, dfm -> 13.000000, df -> 2856.000000, fStat -> 11.586982, aic -> -21209.636294, bic -> -21126.167352, mape -> 106.118428, smape -> 64.262340)



```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | 0.031 | 0.068 | 0.045 | 0.015 | 0.019 |
[info] | rSqBar | 5 | 0.026 | 0.064 | 0.040 | 0.015 | 0.019 |
[info] | sst | 5 | 65472382.329 | 156684138.578 | 94645824.164 | 39425546.436 | 48962871.126 |
[info] | sse | 5 | 61008561.271 | 150517814.952 | 90641793.955 | 38346216.804 | 47622443.854 |
[info] | mse0 | 5 | 106286.692 | 262226.158 | 157912.533 | 66805.256 | 82965.930 |
[info] | rmse | 5 | 326.016 | 512.080 | 390.887 | 79.998 | 99.350 |
[info] | mae | 5 | 134.522 | 150.629 | 140.623 | 6.005 | 7.458 |
[info] | dfm | 5 | 13.000 | 13.000 | 13.000 | 0.000 | 0.000 |
[info] | df | 5 | 2856.000 | 2856.000 | 2856.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 6.972 | 16.074 | 10.280 | 3.732 | 4.635 |
[info] | aic | 5 | -4412.916 | -4127.232 | -4221.811 | 122.389 | 151.995 |
[info] | bic | 5 | -4351.979 | -4066.295 | -4160.875 | 122.389 | 151.995 |
[info] | mape | 5 | 103.255 | 114.642 | 107.370 | 4.494 | 5.581 |
[info] | smape | 5 | 62.595 | 66.148 | 64.796 | 1.397 | 1.735 |
[info] -----
```



Symbolic Lasso Regression:

Multiple R-squared: 0.050100

Adjusted R-squared: 0.045776

**Symbolic Ridge Regression:**

[info] REPORT

[info] modelName mn = SymRidgeRegression.quadratic

[info] hparameter hp = HyperParameter(HashMap(lambda -> (0.01,0.01), factorization -> (Fac\_Cholesky,Fac\_Cholesky)))

[info] features fn = Array(region, accommodation\_type, yearly\_availability, minimum\_nights, number\_of\_reviews, reviews\_per\_month, owned\_hotels, region^2, accommodation\_type^2, yearly\_availability^2, minimum\_nights^2, number\_of\_reviews^2, reviews\_per\_month^2, owned\_hotels^2)


[info] parameter b = VectorD(-93.5564, 376.299, 26.1966, 0.373350, -0.130273, -44.6175, -1.65224, 6.66527, -90.1907, 26.1966, -0.000582615, 0.000987056, 5.30785, 0.0054022)

```
[info] fitMap qof = LinkedHashMap(rSq -> 0.050100, rSqBar -> 0.045442, sst ->
473344895.742509, sse -> 449630543.023248, mse0 -> 156665.694433, rmse -> 395.810175,
mae -> 139.208477, dfm -> 14.000000, df -> 2855.000000, fStat -> 10.755573, aic ->
-21207.636295, bic -> -21118.205285, mape -> 106.117717, smape -> 64.261292)
```

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] |
```

	name	num	min	max	mean	stdev	interval
[info]	rSq	5	0.031	0.068	0.045	0.015	0.019
[info]	rSqBar	5	0.026	0.064	0.040	0.015	0.019
[info]	sst	5	65472382.329	156684138.578	94645824.164	39425546.436	48962871.126
[info]	sse	5	61008752.764	150517860.398	90641807.607	38346144.710	47622354.320
[info]	mse0	5	106287.026	262226.238	157912.557	66805.130	82965.774
[info]	rmse	5	326.017	512.080	390.887	79.997	99.349
[info]	mae	5	134.521	150.628	140.622	6.005	7.458
[info]	dfm	5	14.000	14.000	14.000	0.000	0.000
[info]	df	5	2855.000	2855.000	2855.000	0.000	0.000
[info]	fStat	5	6.471	14.920	9.543	3.464	4.302
[info]	aic	5	-4410.916	-4125.232	-4219.812	122.388	151.995
[info]	bic	5	-4345.627	-4059.943	-4154.522	122.388	151.995
[info]	mape	5	103.255	114.641	107.369	4.494	5.581
[info]	smape	5	62.594	66.147	64.795	1.397	1.735

```
[info] -----
```



Symbolic Lasso Regression:

Multiple R-squared: 0.050100

Adjusted R-squared: 0.045442


## Auto MPG:

### Data Set Information:

This dataset is a slightly modified version of the dataset provided in the StatLib library. In line with the use by Ross Quinlan (1993) in predicting the attribute "mpg", 8 of the original instances were removed because they had unknown values for the "mpg" attribute. The original dataset is available in the file "auto-mpg.data-original".

### Attribute Information:

1. mpg: continuous
2. cylinders: multi-valued discrete
3. displacement: continuous
4. horsepower: continuous
5. weight: continuous
6. acceleration: continuous
7. model year: multi-valued discrete
8. origin: multi-valued discrete
9. car name: string (unique for each instance)



### Scala Methods to Execute:

LasoTest9  
symLassoRegressionTest11  
regressionTest9  
QuadraticTest9  
RidgeRegressionTest9  
symbolicRegressionTest9  
symRidgeRegressionTest11

### Models:

#### Linear Regression:

Multiple R-squared: 0.792574  
Adjusted R-squared: 0.789341

#### Lasso Regression

Multiple R-squared: 0.804643,  
Adjusted R-squared: 0.802112

#### Ridge Regression

Multiple R-squared: 0.804643,  
Adjusted R-squared: 0.801598

#### Quadratic Regression

```
[info] |-----|
[info] | showQofStatTable: Statistical Table for QoF |
[info] |-----|
[info] | name | num | min | max | mean | stdev | interval |
[info] |-----|
[info] | rSq | 5 | 0.888 | 0.923 | 0.903 | 0.013 | 0.016 |
[info] | rSqBar | 5 | 0.888 | 0.923 | 0.903 | 0.013 | 0.016 |
[info] | sst | 5 | 723406.718 | 981947.295 | 847932.772 | 94732.611 | 117649.115 |
[info] | sse | 5 | 64848.888 | 94144.564 | 81567.946 | 11218.803 | 13932.713 |
[info] | mse0 | 5 | 831.396 | 1206.982 | 1045.743 | 143.831 | 178.625 |
[info] | rmse | 5 | 28.834 | 34.742 | 32.274 | 2.265 | 2.813 |
[info] | mae | 5 | 23.021 | 27.148 | 25.222 | 1.835 | 2.279 |
[info] | dfm | 5 | 2.000 | 2.000 | 2.000 | 0.000 | 0.000 |
[info] | df | 5 | 389.000 | 389.000 | 389.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 1547.303 | 2333.118 | 1849.146 | 292.714 | 363.524 |
[info] | aic | 5 | -381.941 | -367.705 | -375.829 | 5.452 | 6.770 |
[info] | bic | 5 | -374.870 | -360.635 | -368.759 | 5.452 | 6.770 |
[info] | mape | 5 | 13.382 | 16.794 | 14.759 | 1.413 | 1.755 |
[info] | smape | 5 | 12.897 | 15.686 | 14.317 | 1.201 | 1.491 |
[info] |-----|
[info] | Total time: 10 s, completed Feb 21, 2022, 10:22:54 PM
```

```

fname = Array(x0, x1, x2, x3)
SUMMARY
Parameters/Coefficients:
  Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
-----
x0      -1.678041       0.015371      -109.167715    0.000000      NA
x1      -0.272654       0.008015      -34.019350    0.000000      3.899141
x2       0.502796       0.000416      1209.082517    0.000000      0.009250
x3      -0.099925       0.004406      -22.677139    0.000000      1.556891
Residual standard error: 5.049825 on 9564.0 degrees of freedom
Multiple R-squared:  0.918517,    Adjusted R-squared:  0.918483
F-statistic: 26952.511083359277 on 4.0 and 9564.0 DF,  p-value: 0.0
-----

REPORT
modelName  mn  = RidgeRegression
hparameter hp  = HyperParameter (HashMap(lambda -> (1.0,0.01), factorization -> (Fac_Cholesky,Fac_Cholesky))
features    fn  = Array(x0, x1, x2, x3)
parameter   b  = VectorD(-1.67804, -0.272654, 0.502796, -0.099924)
fitMap      qof = LinkedHashMap(rSq -> 0.918517, rSqBar -> 0.918483, sst -> 2993123.926550, sse -> 243888.9909
e0 -> 25.487406, rmse -> 5.048505, mae -> 4.019090, dfm -> 4.000000, df -> 9564.000000, fStat -> 26952.511083, aic

```

Multiple R-squared: 0.898704,  
Adjusted R-squared: 0.898183

#### Symbolic Regression

Multiple R-squared: 0.881253  
Adjusted R-squared: 0.874851

#### Ridge Symbolic Regression

Multiple R-squared: 0.889053  
Adjusted R-squared: 0.881150

#### Lasso Symbolic Regression

Multiple R-squared: 0.889053,  
Adjusted R-squared: 0.881474

### Best Model Quadratic Regression Based on R2

#### FoldsCpp

The dataset contains 9568 data points collected from a Combined Cycle Power Plant over 6 years (2006-2011), when the power plant was set to work with full load. Features consist of hourly average ambient variables Temperature (T), Ambient Pressure (AP), Relative Humidity (RH) and Exhaust Vacuum (V) to predict the net hourly electrical energy output (EP) of the plant. A combined cycle power plant (CCPP) is composed of gas turbines (GT), steam turbines (ST) and heat recovery steam generators. In a CCPP, the electricity is generated by gas and steam turbines, which are combined in one cycle, and is transferred from one turbine to another. While the Vacuum is collected from and has effect on the Steam Turbine, the other three of the ambient variables effect the GT performance.

For comparability with our baseline studies, and to allow 5x2 fold statistical tests be carried out, we provide the data shuffled five times. For each shuffling 2-fold CV is carried out and the resulting 10 measurements are used for statistical testing.

```

info] stepRegressionAll: features in/out = ArrayBuffer(5, 3, 7, 1, 2, 9, 24, 25, 8, 16, 4, 23, 22, 12, 6, 15, 21, 7,
0, 19, 9, 18, 10, 13)
info] k = 21, n = 6
info] Run + title
info] x-axis: minX = 0.0, maxX = 20.0
info] y-axis: minY = 69.0, maxY = 89.0
info] Stepwise: rSq =
info] MatrixD(70.6260, 70.4750, 69.7413,
info]      80.2612, 80.1085, 79.7154,
info]      83.3067, 83.1341, 82.8338,
info]      84.6036, 84.4042, 84.0090,
info]      85.5426, 85.3173, 84.9240,
info]      86.1867, 85.9349, 85.5414,
info]      86.2924, 86.0061, 85.5695,
info]      86.6843, 86.3706, 85.7534,
info]      86.8779, 86.5335, 86.0013,
info]      87.0722, 86.6980, 86.1777,
info]      87.3721, 86.9723, 86.3406,
info]      87.4492, 87.0176, 86.1168,
info]      87.5822, 87.1211, 86.1419,
info]      87.6392, 87.1461, 86.0202,
info]      87.9441, 87.4297, 86.3310,
info]      88.1138, 87.5735, 86.4840,
info]      88.2580, 87.6914, 86.5743,
info]      88.2479, 87.7138, 86.5770,
info]      88.2861, 87.7208, 86.5971,
info]      88.3907, 87.7978, 86.5435)

```

### Scala Methods to Execute:

```

regressionTest8
LasoTest8
QuadraticTest8
RidgeRegressionTest8
symbolicRegressionTest11
symRidgeRegressionTest10
symLassoRegressionTest10

```

### Models:

#### Linear Regression:

```

Multiple R-squared: 0.918517
Adjusted R-squared: 0.918491

```

```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
[info] -----
[info] x0      -1.678056      0.015370      -109.174372    0.000000      NA
[info] x1      -0.272647      0.008014      -34.020349    0.000000      3.899141
[info] x2      0.502796      0.000416      1209.146016    0.000000      0.009250
[info] x3      -0.099927      0.004406      -22.678864    0.000000      1.556891
[info] Residual standard error: 5.049561 on 9565.0 degrees of freedom
[info] Multiple R-squared:  0.918517,    Adjusted R-squared:  0.918491
[info] F-statistic: 35940.43894310781 on 3.0 and 9565.0 DF,  p-value: 0.0
[info] -----
[info] REPORT
[info] modelName mn = Regression
[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac_QR,Fac_QR)))
[info] features fn = Array(x0, x1, x2, x3)
[info] parameter b = VectorD(-1.67806, -0.272647, 0.502796, -0.099927)
[info] fitMap qof = LinkedHashMap(rSq -> 0.918517, rSqBar -> 0.918491, sst -> 2993123.926550, sse -> 243888.990904, mse0 -> 25.487406, rmse -> 5.048505, mae -> 4.01
9089, dfm -> 3.000000, df -> 9565.000000, fStat -> 35940.438943, aic -> -29062.916372, bic -> -29034.251236, mape -> NaN, smape -> NaN)
[info]
```

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | 0.910 | 0.934 | 0.917 | 0.009 | 0.012 |
[info] | rSqBar | 5 | 0.910 | 0.934 | 0.917 | 0.009 | 0.012 |
[info] | sst | 5 | 549043.244 | 755436.033 | 597905.395 | 88255.289 | 109604.881 |
[info] | sse | 5 | 47882.308 | 50201.425 | 48834.736 | 974.055 | 1209.686 |
[info] | mse0 | 5 | 25.030 | 26.242 | 25.528 | 0.509 | 0.632 |
[info] | rmse | 5 | 5.003 | 5.123 | 5.052 | 0.050 | 0.062 |
[info] | mae | 5 | 3.996 | 4.052 | 4.022 | 0.020 | 0.025 |
[info] | dfm | 5 | 3.000 | 3.000 | 3.000 | 0.000 | 0.000 |
[info] | df | 5 | 9565.000 | 9565.000 | 9565.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 32240.928 | 44790.023 | 35792.136 | 5089.330 | 6320.476 |
[info] | aic | 5 | -5832.081 | -5786.585 | -5805.270 | 19.109 | 23.731 |
[info] | bic | 5 | -5809.855 | -5764.360 | -5783.044 | 19.109 | 23.731 |
[info] | mape | 5 | 0.879 | 0.889 | NaN | NaN | NaN |
[info] | smape | 5 | 0.878 | 0.889 | NaN | NaN | NaN |
[info] -----
[success] Total time: 11 s, completed Feb 21, 2022, 10:36:34 PM
```

## Lasso Regression

Multiple R-squared: 0.918517

Adjusted R-squared: 0.918491

```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
[info] -----
[info] x0      -1.678056      0.015370      -109.174368    0.000000      NA
[info] x1      -0.272647      0.008014      -34.020352    0.000000      3.899141
[info] x2      0.502796      0.000416      1209.146013    0.000000      0.009250
[info] x3      -0.099927      0.004406      -22.678860    0.000000      1.556891
[info] Residual standard error: 5.049561 on 9565.0 degrees of freedom
[info] Multiple R-squared:  0.918517,    Adjusted R-squared:  0.918491
[info] F-statistic: 35940.438943107714 on 3.0 and 9565.0 DF,  p-value: 0.0
[info] -----
[info] REPORT
[info] modelName mn = LassoRegression
[info] hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01)))
[info] features fn = Array(x0, x1, x2, x3)
[info] parameter b = VectorD(-1.67806, -0.272647, 0.502796, -0.099927)
[info] fitMap qof = LinkedHashMap(rSq -> 0.918517, rSqBar -> 0.918491, sst -> 2993123.926550, sse -> 243888.990904, mse0 -> 25.487406, rmse -> 5.048505, mae -> 4.01
9089, dfm -> 3.000000, df -> 9565.000000, fStat -> 35940.438943, aic -> -29062.916372, bic -> -29034.251236, mape -> NaN, smape -> NaN)
[info]
```

```
[info] -----
[info] | showQofStatTable: Statistical Table for QoF |
[info] -----
[info] | name | num | min | max | mean | stdev | interval |
[info] -----
[info] | rSq | 5 | 0.910 | 0.934 | 0.917 | 0.009 | 0.012 |
[info] | rSqBar | 5 | 0.910 | 0.934 | 0.917 | 0.009 | 0.012 |
[info] | sst | 5 | 549043.244 | 755436.033 | 597905.395 | 88255.289 | 109604.881 |
[info] | sse | 5 | 47882.308 | 50201.425 | 48834.736 | 974.055 | 1209.686 |
[info] | mse0 | 5 | 25.030 | 26.242 | 25.528 | 0.509 | 0.632 |
[info] | rmse | 5 | 5.003 | 5.123 | 5.052 | 0.050 | 0.062 |
[info] | mae | 5 | 3.996 | 4.052 | 4.022 | 0.020 | 0.025 |
[info] | dfm | 5 | 3.000 | 3.000 | 3.000 | 0.000 | 0.000 |
[info] | df | 5 | 9565.000 | 9565.000 | 9565.000 | 0.000 | 0.000 |
[info] | fStat | 5 | 32240.928 | 44790.024 | 35792.136 | 5089.330 | 6320.476 |
[info] | aic | 5 | -5832.081 | -5786.585 | -5805.270 | 19.109 | 23.731 |
[info] | bic | 5 | -5809.855 | -5764.360 | -5783.044 | 19.109 | 23.731 |
[info] | mape | 5 | 0.879 | 0.889 | NaN | NaN | NaN |
[info] | smape | 5 | 0.878 | 0.889 | NaN | NaN | NaN |
[info] -----
```


## Ridge Regression

```
[info] SUMMARY
[info] Parameters/Coefficients:
[info] Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
[info] -----
[info] x0      -1.678041      0.015371      -109.167715    0.000000      NA
[info] x1       -0.272654      0.008015      -34.019350    0.000000      3.899141
[info] x2       0.502796      0.000416      1209.082517    0.000000      0.009250
[info] x3      -0.099925      0.004406      -22.677139    0.000000      1.556891
[info] Residual standard error: 5.049825 on 9564.0 degrees of freedom
[info] Multiple R-squared:  0.918517,    Adjusted R-squared:  0.918483
[info] F-statistic: 26952.511083359277 on 4.0 and 9564.0 DF,  p-value: 0.0
[info] -----
[info] REPORT
[info] modelName mn = RidgeRegression
[info] hparameter hp = HyperParameter (HashMap(lambda -> (1.0,0.01), factorization -> (Fac_Cholesky,Fac_Cholesky)))
[info] features fn = Array(x0, x1, x2, x3)
[info] parameter b = VectorD(-1.67804, -0.272654, 0.502796, -0.099924)
[info] fitMap qof = LinkedHashMap(rSq -> 0.918517, rSqBar -> 0.918483, sst -> 2993123.926550, sse -> 243888.9909
27, mse0 -> 25.487406, rmse -> 5.048505, mae -> 4.019090, dfm -> 4.000000, df -> 9564.000000, fStat -> 26952.511083, aic
-> -29060.916372, bic -> -29025.084952, mape -> NaN, smape -> NaN)
[info] 1.5e+01 0.0e+00 0.0e+00 1.0e+01 0.0e+00 0.0e+00 0.0e+00 0.0e+00 0.0e+00 0.0e+00
```


Multiple R-squared: 0.918517

Adjusted R-squared: 0.918483

## Quadratic Regression

 PolyRegression: y actual, predicted

Reset Plot



Switch max and min value



Multiple R-squared: 0.722751  
Adjusted R-squared: 0.722693

#### Symbolic Regression

Multiple R-squared: 0.940227  
Adjusted R-squared: 0.940170

#### Ridge Symbolic Regression

info]	-----						
info]	showQofStatTable: Statistical Table for QoF						
info]	-----						
info]		name		num		min	
info]						max	
info]						mean	
info]						stdev	
info]						interval	
info]		-----					
info]		rSq		5		0.926	
info]						0.946	
info]		rSqBar		5		0.926	
info]						0.946	
info]		sst		5		549043.244	
info]						755436.033	
info]		sse		5		38125.965	
info]						41044.776	
info]		mse0		5		19.930	
info]						21.456	
info]		rmse		5		4.464	
info]						4.632	
info]		mae		5		3.568	
info]						3.637	
info]		dfm		5		8.000	
info]						8.000	
info]		df		5		9560.000	
info]						9560.000	
info]		fStat		5		14870.169	
info]						20799.177	
info]		aic		5		-5630.054	
info]						-5558.870	
info]		bic		5		-5580.046	
info]						-5508.862	
info]		map		5		0.785	
info]						0.801	
info]						NaN	
info]						NaN	
info]						NaN	
info]		-----					

Multiple R-squared: 0.934456,  
Adjusted R-squared: 0.934401

#### Lasso Symbolic Regression

-----					
SUMMARY					
Parameters/Coefficients:					
Var	Estimate	Std. Error	t value	Pr(> t )	VIF
-----					
x0	1.124533	0.114373	9.832133	0.000000	NA
x1	1.181933	0.062586	18.884876	0.000000	295.495848
x2	0.429808	0.002535	169.545283	0.000000	0.427206
x3	1.092398	0.031154	35.064556	0.000000	96.718200
x4	-0.000052	0.000002	-25.804227	0.000000	760.425537
x5	0.000883	0.000026	34.240191	0.000000	529.005085
x6	-0.000048	0.000001	-31.857766	0.000000	236.374465
x7	-0.000023	0.000001	-30.586820	0.000000	291.846128
Residual standard error: 4.529796 on 9561.0 degrees of freedom					
Multiple R-squared: 0.934456, Adjusted R-squared: 0.934408					
F-statistic: 19472.787284755657 on 7.0 and 9561.0 DF, p-value: 0.0					
-----					

Multiple R-squared: 0.934456  
Adjusted R-squared: 0.934408

#### Best Model Symbolic Regression Based on R2

## USA Housing Dataset:

We have taken this dataset from Kaggle. The main motive is to predict the price of the house. This dataset contains 6 attributes which are - Avg. Area Income, Avg. Area House Age, Avg. Area Number of Rooms, Avg. Area Number of Bedrooms, Area Population, Price.

## Linear Regression:


### SUMMARY

Parameters/Coefficients:					
Var	Estimate	Std. Error	t value	Pr(> t )	VIF
x0	10.237202	0.268622	38.110015	0.000000	NA
x1	51296.068671	2960.620169	17.326123	0.000000	0.740165
x2	-9614.433350	3262.670492	-2.946799	0.003211	0.927345
x3	4082.312867	3132.019425	1.303412	0.192434	1.276879
x4	8.091825	0.326978	24.747299	0.000000	0.900695
Residual standard error: 242106.334210 on 4996.0 degrees of freedom					
Multiple R-squared: 0.531341, Adjusted R-squared: 0.530966					
F-statistic: 1416.0525157356897 on 4.0 and 4996.0 DF, p-value: 0.0					


### [info] REPORT

```
[info] modelName mn = Regression
[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac_QR, Fac_QR)))
[info] features fn = Array(x0, x1, x2, x3, x4)
[info] parameter b = VectorD(10.2372, 51296.1, -9614.43, 4082.31, 8.0918)
[info] fitMap qof = LinkedHashMap(rSq -> 0.531341, rSqBar -> 0.530966, sst -> 624853298447502.600000, sse -> 292842923414396.300000, mse0 -> 58556873308.217620, rmse -> 241985.274982, mae -> 193470.071203, dfm -> 4.000000, df -> 4996.000000, fStat -> 1416.052516, aic -> -69081.710487, bic -> -69049.123521, mape -> NaN, smape -> NaN)
```


R<sup>2</sup> vs n for Regression with Forward



R<sup>2</sup> vs n for Regression with Backward



R<sup>2</sup> vs n for Regression with Stepwise




Lasso Regression:

SUMMARY


Parameters/Coefficients:					
Var	Estimate	Std. Error	t value	Pr(> t )	VIF
x0	10.237202	0.268622	38.110015	0.000000	NA
x1	51296.068669	2960.620169	17.326123	0.000000	0.740165
x2	-9614.433347	3262.670492	-2.946799	0.003211	0.927345
x3	4082.312865	3132.019425	1.303412	0.192434	1.276879
x4	8.091825	0.326978	24.747299	0.000000	0.900695
Residual standard error: 242106.334210 on 4996.0 degrees of freedom					
Multiple R-squared: 0.531341, Adjusted R-squared: 0.530966					
F-statistic: 1416.0525157356935 on 4.0 and 4996.0 DF, p-value: 0.0					

```
[info] REPORT
[info]   modelName mn = LassoRegression
[info]   hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01)))
[info]   features   fn = Array(x0, x1, x2, x3, x4)
[info]   parameter  b = VectorD(10.2372, 51296.1, -9614.43, 4082.31, 8.0918)
[info]   fitMap     qof = LinkedHashMap(rSq -> 0.531341, rSqBar -> 0.530966, sst -> 624853298447502.600000, sse -> 292
342923414395.900000, mse0 -> 58556873308.217530, rmse -> 241985.274982, mae -> 193470.071204, dfm -> 4.000000, df -> 499
6.000000, fStat -> 1416.052516, aic -> -69081.710487, bic -> -69049.123521, mape -> NaN, smape -> NaN)
```




R<sup>2</sup> vs n for Regression with Forward

— □ ×




R<sup>2</sup> vs n for Regression with Backward

— □ ×



R<sup>2</sup> vs n for Regression with Stepwise



## Ridge Regression:

### SUMMARY

#### Parameters/Coefficients:

Var	Estimate	Std. Error	t value	Pr(> t )	VIF
x0	10.237202	0.268622	38.110015	0.000000	NA
x1	51296.068671	2960.620169	17.326123	0.000000	0.740165
x2	-9614.433350	3262.670492	-2.946799	0.003211	0.927345
x3	4082.312867	3132.019425	1.303412	0.192434	1.276879
x4	8.091825	0.326978	24.747299	0.000000	0.900695

Residual standard error: 242106.334210 on 4996.0 degrees of freedom


Multiple R-squared: 0.531341, Adjusted R-squared: 0.530966

F-statistic: 1416.0525157356897 on 4.0 and 4996.0 DF, p-value: 0.0


### [info] REPORT

```
[info] modelName mn = Regression
[info] hparameter hp = HyperParameter (HashMap(factorization -> (Fac_QR,Fac_QR)))
[info] features fn = Array(x0, x1, x2, x3, x4)
[info] parameter b = VectorD(10.2372, 51296.1, -9614.43, 4082.31, 8.0918)
[info] fitMap qof = LinkedHashMap(rSq -> 0.531341, rSqBar -> 0.530966, sst -> 624853298447502.600000, sse -> 842923414396.300000, mse0 -> 58556873308.217620, rmse -> 241985.274982, mae -> 193470.071203, dfm -> 4.000000, df -> 6.000000, fStat -> 1416.052516, aic -> -69081.710487, bic -> -69049.123521, mape -> NaN, smape -> NaN)
[info]
```


R<sup>2</sup> vs n for RidgeRegression with Forward



R<sup>2</sup> vs n for RidgeRegression with Backward



R<sup>2</sup> vs n for RidgeRegression with Stepwise





## Quadratic Regression:


```
SUMMARY
Parameters/Coefficients:
Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
-----
x0       5.136085        0.375942        13.661891     0.000000      NA
x1       0.000025        0.000011        2.235661     0.025374     70.737970
x2      -0.000000        0.000000       -2.203290     0.027574     70.737970
Residual standard error: 0.994654 on 4998.0 degrees of freedom
Multiple R-squared:  0.001003,    Adjusted R-squared:  0.000603
F-statistic: 2.508735121336463 on 2.0 and 4998.0 DF,  p-value: 0.0
```

```
[info] REPORT
[info]   modelName mn = PolyRegression
[info]   hparameter hp = HyperParameter (HashMap(factorization -> (Fac_Cholesky,Fac_Cholesky)))
[info]   features   fn = Array(x0, x1, x2)
[info]   parameter  b = VectorD(5.13608, 2.47170e-05, -1.77501e-1)
[info]   fitMap     qof = LinkedHashMap(rSq -> 0.001003, rSqBar -> 0.000603, sst -> 4949.663837, sse -> 4944.699874, mse0 -> 0.988742, rmse -> 0.994355, mae -> 0.795422, dfm -> 2.000000, df -> 4998.000000, fStat -> 2.508735, aic -> -7061.801889, bic -> -7042.249709, mape -> Infinity, smape -> 13.516651)
```






R<sup>2</sup> vs n for Regression with Forward



R<sup>2</sup> vs n for Regression with Backward







Symbolic Regression:

SUMMARY

Parameters/Coefficients:					
Var	Estimate	Std. Error	t value	Pr(> t )	VIF
x0	0.748740	0.968482	0.773107	0.439459	NA
x1	-59820.197647	11031.304218	-5.422767	0.000000	53.835975
x2	-116041.009552	11497.329285	-10.092866	0.000000	60.331523
x3	-50652.522718	16740.784098	-3.025696	0.002481	191.120743
x4	-1.247255	1.482206	-0.841485	0.400076	96.964524
x5	1.443778	0.125642	11.491168	0.000000	61.922070
x6	1.169153	0.138727	8.427694	0.000000	89.127736
x7	15293.731548	1534.481878	9.966707	0.000000	89.547897
x8	0.131378	0.127026	1.034256	0.301016	67.160827
x9	-1695.790808	1379.002077	-1.229723	0.218801	61.331687
x10	7681.708699	1356.387824	5.663357	0.000000	104.384962
x11	0.000090	0.000014	6.657232	0.000000	50.560730
x12	0.681448	0.141904	4.802189	0.000002	43.444824
x13	0.901487	0.160334	5.622567	0.000000	72.323231
x14	-0.043388	0.137809	-0.314839	0.752884	30.348317
Residual standard error: 105773.876897 on 4986.0 degrees of freedom					
Multiple R-squared: 0.910725, Adjusted R-squared: 0.910474					
F-statistic: 3633.1248885464265 on 14.0 and 4986.0 DF, p-value: 0.0					

```
[info] REPORT
[info]   modelName mn = SymbolicRegressionX
[info]   hparameter hp = HyperParameter(HashMap(factorization -> (Fac_QR,Fac_QR)))
[info]   features fn = Array(Avg. Area Income, Avg. Area House Age, Avg. Area Number of
Bedrooms, Area Population, Avg. Area House Age_Avg. Area Income, Avg. Area N
umber of Rooms_Avg. Area House Age, Avg. Area Number of Bedrooms_Avg. Area Income, Avg. Area Number of Bedrooms_Avg. Are
a House Age, Avg. Area Number of Bedrooms_Avg. Area Number of Rooms, Area Population_Avg. Area Income, Area Population_A
vg. Area House Age, Area Population_Avg. Area Number of Rooms, Area Population_Avg. Area Number of Bedrooms)
[info]   parameter b = VectorD(0.748740, -59820.2, -116041, -50652.5, -1.24725, 1.44378,
1.16915, 15293.7, 0.131378, -1695.79, 7681.71, 9.00561e-05, 0.681448, 0.901487
, -0.043387)
[info]   fitMap qof = LinkedHashMap(rSq -> 0.910725, rSqBar -> 0.910474, sst -> 624853298447502.600000, sse -> 557
83931586116.445000, mse0 -> 11154555406.142061, rmse -> 105615.128680, mae -> 84679.343701, dfm -> 14.000000, df -> 4986
.000000, fStat -> 3633.124889, aic -> -64915.463682, bic -> -64817.702784, mape -> NaN, smape -> NaN)
[info]
```







### Symbolic Lasso Regression:

SUMMARY					
Parameters/Coefficients:					
Var	Estimate	Std. Error	t value	Pr(> t )	VIF
x0	6.049819	0.511927	11.817730	0.000000	NA
x1	-652.522107	5866.278979	-0.111233	0.911432	13.861542
x2	-56546.883536	6073.730779	-9.310074	0.000000	15.329525
x3	-40405.242861	8796.889354	-4.593128	0.000004	48.048714
x4	2.957676	0.787922	3.753766	0.000174	24.947663
x5	0.232008	0.014462	16.042840	0.000000	51.193915
x6	-0.030833	0.023823	-1.294232	0.195585	96.059965
x7	0.051243	0.016452	3.114658	0.001842	78.264382
x8	719.782992	196.628326	3.660627	0.000252	86.040597
x9	0.000015	0.000002	6.831178	0.000000	54.746866
x10	0.000005	0.000002	2.253994	0.024197	86.659911
x11	0.083322	0.026699	3.120784	0.001804	83.595858
x12	0.000000	0.000004	0.022806	0.981805	121.812318
x13	-0.132079	0.045437	-2.906840	0.003651	126.901650
x14	0.104313	0.030581	3.411038	0.000647	95.514297
Residual standard error: 110852.319815 on 4986.0 degrees of freedom					
Multiple R-squared: 0.901946, Adjusted R-squared: 0.901671					
F-statistic: 3275.9794038948385 on 14.0 and 4986.0 DF, p-value: 0.0					


```
[info] REPORT
[info]   modelName mn = SymLassoRegressionXX
[info]   hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01)))
[info]   features fn = Array(Avg. Area Income, Avg. Area House Age, Avg. Area Number of Rooms, Avg. Area Number of Bedrooms, Area Population, Avg. Area Number of Rooms_Avg. Area House Age_Avg. Area Income, Avg. Area Number of Bedrooms_Avg. Area Number of Rooms_Avg. Area Income, Avg. Area Number of Bedrooms_Avg. Area Number of Rooms_Avg. Area House Age, Area Population_Avg. Area House Age_Avg. Area Income, Area Population_Avg. Area Number of Rooms_Avg. Area Income, Area Population_Avg. Area Number of Bedrooms_Avg. Area House Age, Area Population_Avg. Area Number of Bedrooms_Avg. Area Income, Area Population_Avg. Area Number of Bedrooms_Avg. Area Number of Rooms)
[info]   parameter b = VectorD(6.04982, -652.522, -56546.9, -40405.2, 2.95768, 0.232008, -0.0308328, 0.0512426, 719.783, 1.49384e-05, 5.41563e-06, 0.0833218, 8.89805e-08, -0.132079, 0.10431)
[info]   fitMap qof = LinkedHashMap(rSq -> 0.901946, rSqBar -> 0.901671, sst -> 624853298447502.600000, sse -> 61259148726467.500000, mse0 -> 12251379469.399620, rmse -> 110685.949738, mae -> 88342.124664, dfm -> 14.000000, df -> 4986.000000, fStat -> 3275.979404, aic -> -65149.987003, bic -> -65052.226105, mape -> NaN, smape -> NaN)
[info]
```




R<sup>2</sup> vs n for Symbolic Lasso Regression with Forward



R<sup>2</sup> vs n for Symbolic Lasso Regression with Backward



R<sup>2</sup> vs n for Symbolic Lasso Regression with Stepwise




## Symbolic Ridge Regression:

```
SUMMARY
Parameters/Coefficients:
  Var      Estimate      Std. Error      t value      Pr(>|t|)      VIF
-----
x0      6.049747      0.511979      11.816404      0.000000      NA
x1     -653.277645     5866.867342     -0.111350      0.911339     13.861542
x2    -56545.921993    6074.339949     -9.308982      0.000000     15.329525
x3   -40403.726429    8797.771646     -4.592495      0.000004     48.048714
x4     2.957598      0.788001      3.753290      0.000175     24.947663
x5     0.232008      0.014463     16.041286      0.000000     51.193915
x6    -0.030831      0.023826     -1.294025      0.195657     96.059965
x7     0.051240      0.016454      3.114209      0.001844     78.264382
x8     719.759222     196.648047      3.660139      0.000252     86.040597
x9     0.000015      0.000002      6.830608      0.000000     54.746866
x10    0.000005      0.000002      2.253767      0.024211     86.659911
x11    0.083321      0.026702      3.120459      0.001806     83.595858
x12    0.000000      0.000004      0.022802      0.981808     121.812318
x13   -0.132079      0.045442     -2.906555      0.003654     126.901650
x14    0.104313      0.030584      3.410690      0.000648     95.514297
Residual standard error: 110863.437846 on 4985.0 degrees of freedom
Multiple R-squared: 0.901946, Adjusted R-squared: 0.901651
F-statistic: 3056.967543696419 on 15.0 and 4985.0 DF, p-value: 0.0
```


```
[info] REPORT
[info]   modelName mn = SymRidgeRegressionXX
[info]   hparameter hp = HyperParameter (HashMap(lambda -> (0.01,0.01), factorization -> (Fac_Cholesky,Fac_Cholesky)))
[info]   features fn = Array(Avg. Area Income, Avg. Area House Age, Avg. Area Number of Rooms, Avg. Area Number of Bedrooms, Area Population, Avg. Area Number of Rooms_Avg. Area House Age_Avg. Area Income, Avg. Area Number of Bedrooms_Avg. Area House Age_Avg. Area Income, Avg. Area Number of Bedrooms_Avg. Area Number of Rooms_Avg. Area House Age, Area Population_Avg. Area House Age_Avg. Area Income, Area Population_Avg. Area Number of Rooms_Avg. Area Income, Area Population_Avg. Area Number of Bedrooms_Avg. Area House Age, Area Population_Avg. Area Number of Bedrooms_Avg. Area Income, Area Population_Avg. Area Number of Bedrooms_Avg. Area House Age, Area Population_Avg. Area Number of Bedrooms_Avg. Area Number of Rooms)
[info]   parameter b = VectorD(6.04975, -653.278, -56545.9, -40403.7, 2.95760, 0.232008, -0.0308309, 0.0512403, 719.759, 1.49387e-05, 5.41563e-06, 0.0833215, 8.89750e-08, -0.132079, 0.10431)
[info]   fitMap qof = LinkedHashMap(rSq -> 0.901946, rSqBar -> 0.901651, sst -> 624853298447502.600000, sse -> 61269148727619.380000, mse0 -> 12251379469.629951, rmse -> 110685.949739, mae -> 88342.129691, dfm -> 15.000000, df -> 4985.000000, fStat -> 3056.967544, aic -> -65147.987003, bic -> -65043.708712, mape -> NaN, smape -> NaN)
```

SymRidgeRegressionXX: y actual, predicted


Reset Plot




R<sup>2</sup> vs n for Symbolic Ridge Regression with Forward



R<sup>2</sup> vs n for Symbolic Ridge Regression with Backward



R<sup>2</sup> vs n for Symbolic Ridge Regression with Stepwise



### Scala Methods to Execute:

```
LassoRegressionTest100
symLassoRegressionTest10
LinearRegression100
QuadraticTest100
RidgeRegressionTest9
symbolicRegressionTest100
symRidgeRegressionTest10
```



## Models:

Linear Regression:

Multiple R-squared: 0.531341

Adjusted R-squared: 0.530966

Lasso Regression

Multiple R-squared: 0.530966

Adjusted R-squared: 0.531341

Ridge Regression

Multiple R-squared: 0.531341

Adjusted R-squared: 0.530966

Quadratic Regression

Multiple R-squared: 0.001003

Adjusted R-squared: 0.0.000603

Symbolic Regression

Multiple R-squared: 0.910725

Adjusted R-squared: 0.910474

Ridge Symbolic Regression

Multiple R-squared: 0.901946

Adjusted R-squared: 0.901651

Lasso Symbolic Regression

Multiple R-squared: 0.901946

Adjusted R-squared: 0.901671

**Best Model** Symbolic Lasso Regression Regression & Symbolic Ridge Regression